



BITCOIN AND BLOCKCHAIN TECHNOLOGY WITH FOCUS ON FPGA MINERS

**Introduction
Cryptography And Cryptocurrencies
Bitcoin Mining
FPGA
Conclusion**



**MOHAMMAD NIKNAM
951843189
2021**

Table of Contents

1.INTRODUCTION.....	3
2.CRYPTOGRAPHY AND CRYPTOCURRENCIES	4
2.1-Cryptographic Hash Functions.....	4
Property 1: Collision Resistance.....	4
Property 2: Hiding - Preimage resistance (one wayness)	6
Property 3: Puzzle Friendliness - Target Collision Resistance	7
SHA-256.....	7
2.2-HASH POINTERS AND DATA STRUCTURES.....	10
Hash pointer.....	10
Block chain	10
Merkle tree	11
2.3-DIGITAL SIGNATURES.....	14
Digital signature scheme.....	14
ECDSA.....	15
Public Keys as Identities.....	16
3-BITCOIN MINING.....	18
3.1-The task of Bitcoin miners	18
Finding a valid block.....	19
Difficulty	20
3.2-Mining Hardware	24
A closer look at SHA-256.....	24
CPU mining.....	25
GPU mining	25
FPGA mining.....	27
ASIC mining	28
4-FPGA	29
4.1-Introduction	29
Requirement of FPGAs.....	29
4.2-Architecture and Structure of FPGA	31
Internal Architecture of LUT	32
Routing Architecture - switch matrix.....	33
Additional Elements In Contemporary FPGA Architectures	33
4.3-FPGA Programming.....	35
HDL & Synthesize	35
Translate process	35
Map process.....	35
Place and Route	35
Device Programming - Bitstream Generation.....	36
4.4-FPGAs Vs ASICs	37
4.5-FPGA As Bitcoin Miner.....	38

5-CONCLUSION	39
Professional mining.....	39
Mining hardware evolution: Similarities to gold mining	39

1. INTRODUCTION

Bitcoin, "A Peer-to-Peer Electronic Cash System"¹, is unlike anything the world has seen before. By providing fast, inexpensive, international money transfer, it has the potential to revolutionize both the modern-day concept of money and commerce.

Bitcoin is a peer-to-peer² digital currency based on public key cryptography. It started as a free software project and a paper published by Satoshi Nakamoto³ in 2009. Nakamoto, designed a system of online value transfer that supports a promising Internet currency. He introduced bitcoin as a version of electronic cash that would allow payments to be sent from one party to another without going through a financial institution.

Traditionally, financial institutions, such as banks, are trusted to store and protect a customer's currency. The bank will handle the transfer of money between its customers and clients, but there are several disadvantages in this trust-based model. Electronic transfers between banks can be costly since there is usually a transaction fee, they can be slow taking several days to complete, and transfers cannot be made anonymously and completely non-reversible transactions are not really possible, since financial institutions cannot avoid mediating disputes. The cost of mediation increases transaction costs. But no mechanism exists to make payments over a communications channel without a trusted party.

What is needed is an electronic payment system based on cryptographic proof instead of trust. Bitcoin is a system of owning and transferring currency that omits these trusted third parties, allowing any two willing parties to transact directly with each other without the need for a trusted third party. And instead, Satoshi propose a solution to the double-spending⁴ problem using a peer-to-peer distributed Timestamp server to generate computational proof of the chronological order of transactions.

Bitcoin is made possible by a combination of software and network technologies. A program called the Bitcoin client simultaneously manages and helps you spend bitcoins. This program maintains a long ledger called the blockchain that holds every transaction confirmed by the Bitcoin network.

The Bitcoin network, consisting of thousands of machines running the Bitcoin software, has two main tasks to accomplish. One is relaying transaction information and the second is verifying those transactions to ensure the same bitcoins cannot be spent twice.

The first task is accomplished easily due to the fact that the Bitcoin network is operated as a peer-to-peer network. After all, sharing data is easy. By operating on many nodes across the globe, the network ensures it will operate as long as it provides a useful service, and the system is secure as long as honest nodes collectively control more CPU power than any cooperating group of attacker nodes. The second task is a bit more complicated and is solved through what I consider to be Bitcoin's key innovation. This development, a process called mining, is carried out by computers running mining software.

¹ This phrase is come from title of original bitcoin whitepaper by Nakamoto in 2009.

² P2P - Peer-to-peer refers to systems that work like an organized collective by allowing each individual to interact directly with the others. In the case of Bitcoin, the network is built in such a way that each user is broadcasting the transactions of other users. And, crucially, no bank is required as a third party.

³ Satoshi Nakamoto – the creator of Bitcoin and the author of the original Bitcoin whitepaper and code. His real identity is unknown to the world.

⁴ Double Spend - If a malicious user tries to spend their bitcoins with two different recipients at the same time, this is double spending. Bitcoin mining and the block chain are there to create a consensus on the network about which of the two transactions will confirm and be considered valid.

2.CRYPTOGRAPHY AND CRYPTOCURRENCIES

Cryptography is a deep academic research field using many advanced mathematical techniques that are notoriously subtle and complicated. Fortunately, Bitcoin relies on only a handful of relatively simple and well-known cryptographic constructions. In this section, we specifically study cryptographic hashes and digital signatures, two primitives that prove to be useful for building cryptocurrencies.

2.1-Cryptographic Hash Functions

In this section, we've talked about hash functions, cryptographic hash functions with special properties, applications of those properties, and a specific hash function that we use in Bitcoin.

The first cryptographic primitive that we need to understand is a *cryptographic hash function*. A *hash function* is a mathematical function with the following three properties:

- Its input can be any string of any size.
- It produces a fixed-sized output. For the purpose of making the discussion in this chapter concrete, we will assume a 256-bit output size. However, our discussion holds true for any output size, as long as it is sufficiently large.
- It is efficiently computable. Intuitively this means that for a given input string, you can figure out what the output of the hash function is in a reasonable amount of time. More technically, computing the hash of an n -bit string should have a running time that is $O(n)$.

These properties define a general hash function, one that could be used to build a data structure, such as a hash table. We're going to focus exclusively on *cryptographic* hash functions. For a hash function to be cryptographically secure, we require that it has the following three additional properties: (1) collision resistance, (2) hiding, and (3) puzzle friendliness.

Property 1: Collision Resistance

The first property that we need from a cryptographic hash function is that it's collision-resistant. A collision occurs when two distinct inputs produce the same output. A hash function $H(\cdot)$ is collision-resistant if nobody can find a collision.

Formally:

Collision-resistance: A hash function H is said to be collision resistant if it is infeasible to find two values, x and y , such that $[x \neq y]$, yet $[H(x) = H(y)]$.

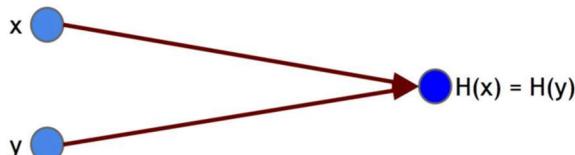


Figure 1. A hash collision. x and y are distinct values, yet when input into hash function H , they produce the same output.

Notice that we said “nobody can find” a collision, but we did not say that no collisions exist. Actually, collisions exist for any hash function, and we can prove this by a simple counting argument. The input space to the hash function contains all strings of all lengths, yet the output space contains only strings of a specific fixed length. Because the input space is larger than the output space (indeed, the input space is infinite, while the output space is finite), there must be input strings that map to the same

output string. In fact, there will be some outputs to which an infinite number of possible inputs will map.

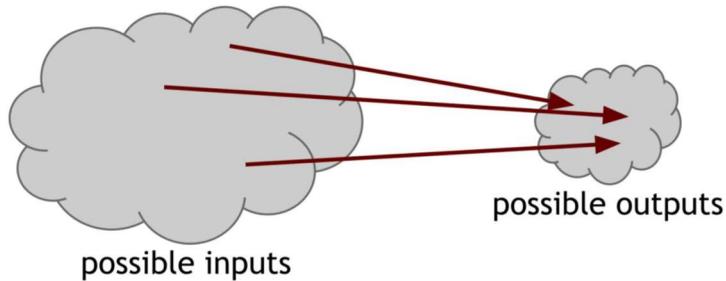


Figure 2. Inevitability of collisions. Because the number of inputs exceeds the number of outputs, we are guaranteed that there must be at least one output to which the hash function maps more than one input.

Now, to make things even worse, we said that it has to be impossible to find a collision. Yet, there are methods that are guaranteed to find a collision. Consider the following simple method for finding a collision for a hash function with a 256-bit output size: pick $2^{256} + 1$ distinct values, compute the hashes of each of them, and check if there are any two outputs are equal. Since we picked more inputs than possible outputs, some pair of them must collide when you apply the hash function. The method above is guaranteed to find a collision. But if we pick random inputs and compute the hash values, we'll find a collision with high probability long before examining $2^{256} + 1$ inputs. In fact, if we randomly choose just $2^{130} + 1$ inputs, it turns out there's a 99.8% chance that at least two of them are going to collide.

This collision-detection algorithm works for every hash function. But, of course, the problem is that it takes a very long time to do. For a hash function with a 256-bit output, you would have to compute the hash function $2^{256} + 1$ times in the worst case, and about 2^{128} times on average. That's of course an astronomically large number—if a computer calculates 10,000 hashes per second, it would take more than one octillion (10^{27}) years to calculate 2^{128} hashes! For another way of thinking about this, we can say that if every computer ever made by humanity had been computing since the beginning of the universe, the odds that they would have found a collision by now are still infinitesimally small. So small that it's far less than the odds that the Earth will be destroyed by a giant meteor in the next two seconds.

We have thus found a general but impractical algorithm to find a collision for *any* hash function. A more difficult question is: Is there some other method that could be used on a particular hash function to find a collision? In other words, although the generic collision detection algorithm is not feasible to use, there may be some other algorithm that can efficiently find a collision for a specific hash function. Yet for other hash functions, we don't know if such methods exist. We suspect that they are collision resistant. However, there are no hash functions *proven* to be collision-resistant. The cryptographic hash functions that we rely on in practice are just functions for which people have tried really, really hard to find collisions and haven't yet succeeded. In some cases, such as the old MD5 hash function, collisions were eventually found after years of work, leading the function to be deprecated and phased out of practical use. And so, we choose to believe that those are collision resistant.

Application: Message digests

If we can assume that we have a hash function that is Collision Resistance, then we can use that hash function as message digest. That means if we know that x and y have the same hash, then it's safe to assume that x and y are the same. Because of Collision Resistance Property, since there's not a collision that we know of, then knowing the hashes are the same, we can assume that the values are the same.

And this let us use the hash as a kind of message digest. Suppose, for example, that we had a file, a really big file. And we wanted to be able to recognize later whether another file was the same as the file we saw the first time, right? So, one way to do that would be to save the whole big file. And then when we saw another file later, just compare them. But because we have hashes that we believe are collision free, it's more efficient to just remember the hash of the original file. Then if someone shows us a new file, and claims that it's the same, we can compute the hash of that new file and compare the hashes. If the hashes are the same, then we conclude that the files must have been the same. And that gives us a very efficient way to remember things we've seen before and recognize them again. And, of course, this is useful because the hash is small, it's only 256 bits, while the original file might be really big. So, hash is useful as a message digest.

Property 2: Hiding - Preimage resistance (one wayness)

The second property that we want from our hash functions is that it's hiding. The hiding property asserts that if we're given the output of the hash function $y = H(x)$, there's no feasible way to figure out what the input, x , was.

در شرایطی که تعداد حالت های ممکن ورودی X محدود و مشخص باشد (در واقع قابل حدس باشد)، هر کسی میتواند ورودی های محتمل را تک تک، هش کند و با خروجی مورد نظر مقایسه کند و بدین ترتیب به ورودی X که خروجی $H(x) = y$ مورد نظر را نتیجه داده، پی ببرد. پس در چنین شرایطی، تابع هش ما عملاً ورودی را مخفی نگرده است.

In order to be able to achieve the hiding property, it needs to be the case that there's no value of x which is particularly likely. That is, x has to be chosen from a set that's, in some sense, very spread out. If x is chosen from such a set, this method of trying a few values of x that are especially likely will not work. Also, we want method that we can hide even an input that's not spread out. We can now be slightly more precise about what we mean by hiding (the double vertical bar \parallel denotes concatenation).

Hiding: A hash function H is hiding if: when a secret value r is chosen from a probability distribution that has *high min-entropy*, then given $H(r \parallel x)$ it is infeasible to find x .

High min-entropy means that the distribution is very spread out, and it's not predictable, so that no particular value is chosen with more than negligible probability. for a concrete example, if r is chosen uniformly from among all of the strings that are 256 bits long, then any particular string was chosen with probability $1/2^{256}$, which is an infinitesimally small value. So, as long as r was chosen that way, then the hash of r concatenated with x is going to hide x . And that's the hiding property that the hash function will be deemed to have.

۲: مقداری تصادفی و مخفی است که در رمزگاری به آن Key یا Nonce گفته میشود. عبارت Nonce به این موضوع اشاره دارد که مقدار آن فقط یکبار بکار میرود و در هر مرتبه برای مخفی سازی ورودی X ، با یک مقدار تصادفی جدید ترکیب شده و سپس هش میشود.

کاربرد:

کاربرد دو ویژگی، "برخورد ناپذیری" و "مخفی سازی" برای تابع هش، به کمک مثال بدین صورت توضیح داده میشود: فرض کنید: پیامی را بروی نامه ای نوشته و آن را در پاکتی مهر و موم شده قرار داده ایم. در اینجا منظور از پیام، X و منظور از پاکت، تابع هش، و استفاده از Nonce، به مهر و موم کردن تشبيه شده است. حال در نظر بگیرید که ما این نامه یا پیام (X) را بصورت مهر و موم شده (ترکیب X و Nonce) درون پاکت قرار داده ایم (هش کرده ایم) و آن را بروی میزی در انتظار همگان قرار دادیم، حال، هیچکس جز ما، که آن را مهر و موم کرده ایم (از پیام و Nonce اطلاع داریم)، با نگاه کردن به پاکت (خروجی هش شده) نمیتواند از متن پیام اطلاع پیدا کند (به مقدار X پی ببرد) که این همان ویژگی "مخفی سازی" است. همچنین همانطور که مشخص است، پس از مهر و موم کردن پاکت، حتی اگر مثلاً ما تغییر نظر بدیم، نمیتوانیم پیام درون پاکت را تغییر دهیم و یا آن را با پیامی دیگر جایگزین کنیم (عملانه نمیتوان X دیگری را یافت، که پس از ترکیب با Nonce و هش کردن ما را به همان خروجی قبل برساند) و این همان ویژگی "برخورد ناپذیری" را بیان میکند.

Property 3: Puzzle Friendliness - Target Collision Resistance

این ویژگی یکمقدار پیچیده‌تر است. اگر شخصی-خروجی خاصی مانند \mathcal{H} از تابع هش را مدنظر داشته باشد یا به عنوان هدف در نظر بگیرد، و اگر قسمتی از ورودی بطرز مناسبی تصادف انتخاب شده باشد، آنگاه پیدا کردن مقداری دیگر در ورودی که آن را به خروجی مدنظر برساند، بسیار مشکل و غیر عملی است.

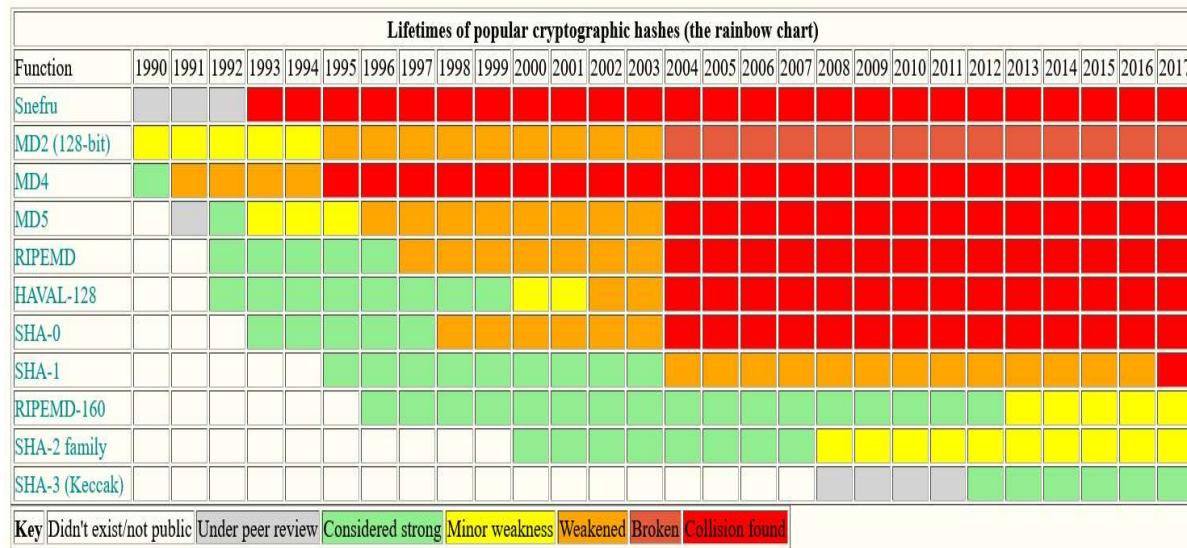
Puzzle friendliness: A hash function H is said to be puzzle friendly if for every possible n -bit output value y , if k is chosen from a distribution with high min-entropy, then it is infeasible to find x such that $H(k \parallel x) = y$ in time significantly less than 2^n .

ویژگی Puzzle Friendly را میتوان اینگونه نیز توضیح داد که یک پازل خوب نباید به سرعت حل شود و مسیر میانبری برای پیدا کردن قطعات نباشد و هیچ راهی بجز جستجو تصادف تک تک قطعات برای یافتن قطعه مناسب نباشد. به همین ترتیب در تابع هش مناسب، چک کردن ورودی ها برای رسیدن به خروجی خاص باید زمانبر باشد و همچنین نتوان هیچ استراتژی و یا الگوریتم میانبر دیگری که از چک کردن تصادفی تک تک ورودی ها بهینه تر باشد، یافت. ونتیجتا در شرایطی که تعداد ورودی ها بسیار زیاد است عملا نتوان ورودی X که خروجی خاص Y را نتیجه دهد پیدا کرد.

در مورد رابطه ویژگی اول و سوم می توان گفت در صورتی که خاصیت اول صدق کند، خاصیت سوم نیز صادق است. توجه شود که ویژگی سوم، با ویژگی دوم گفته شده، یعنی "برخورد ناپذیری"، تفاوت دارد. اینکه بتوان برای یک ورودی و خروجی خاص مدنظر مانند X و $(X)H$ ، ورودی دیگری X' یافت که هش آن $(X')H = H(X)$ باشد، کافیست تا ویژگی سوم نقض شود. در ویژگی اول گفته میشود که که تابع هش باید به گونه ای باشد که عملانتوان هیچ جفت ورودی و خروجی تصادفی را یافت که ورودی ها یکسان نباشند اما خروجی یکسان نتیجه داده باشند.

در تعیین کاربردی مفهوم ویژگی سوم؛ تابع هش با n بیت خروجی، که هر یک از 2^n مقدار را می تواند داشته باشد، را در نظر بگیرید، برای حل این پازل نیاز به پیدا کردن ورودی داریم که خروجی را در مجموعه \mathcal{Y} می توان آن را انتخاب کرد و در نظر گرفت) که عملاً بسیار کوچکتر از مجموعه کل خروجی ها است، قرار دهد. اندازه مجموعه \mathcal{Y} تعیین کننده میزان سختی پازل ما است اگر مجموعه \mathcal{Y} شامل همه رشته های n بیتی باشد (یعنی همه خروجی ها قابل قبول اند) آنگاه حل پازل ساده و بدیهی است و در حالی که مجموعه \mathcal{Y} فقط شامل یک مقدار باشد (یعنی فقط یک خروجی قابل قبول است) است آنگاه حل پازل در سخت ترین حالت می باشد. کاربرد این مفهوم و قابلیت تعیین سختی پازل در مباحث ماینینگ نقش مهمی را دارد که در ادامه خواهیم دید.

SHA-256



2017-1990 Figure 3 دوره زمانی بکارگیری رایج‌ترین توابع هش و وضعیت آن‌ها

Many hash functions exist, but this is the one Bitcoin uses primarily, and it's a pretty good one to use. It's called *SHA-256* from *SHA-2 family*. Recall that we require that our hash functions work on inputs of arbitrary length. Luckily, as long as we can build a hash function that works on fixed-length inputs, there's a generic method to convert it into a hash function that works on arbitrary length inputs. It's called the *Merkle-Damgard transform*. SHA-256 is one of a number of commonly used hash functions

that make use of this method. In common terminology, the underlying fixed-length collision-resistant hash function is called the **compression function**. It has been proven that if the underlying compression function is collision resistant, then the overall hash function is collision resistant as well. The Merkle-Damgård transform is quite simple. Say the compression function takes inputs of length m and produces an output of a smaller length n . The input to the hash function, which can be of any size, is divided into **blocks** of length $m-n$. The construction works as follows: pass each block together with the output of the previous block into the compression function. Notice that input length will then be $(m-n) + n = m$, which is the input length to the compression function. For the first block, to which there is no previous block output, we instead use an **Initialization Vector (IV)**. This number is reused for every call to the hash function, and in practice you can just look it up in a standards document. The last block's output is the result that you return.

SHA-256 uses a compression function that takes 768-bit input and produces 256-bit outputs. The block size is 512 bits. See Figure 4 for a graphical depiction of how SHA-256 works.

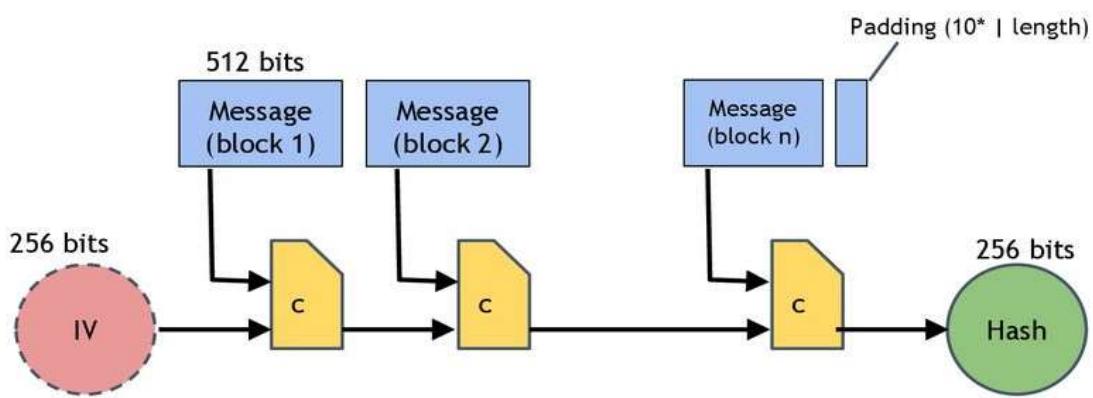


Figure 4: SHA-256 Hash Function (simplified). SHA-256 uses the Merkle-Damgård transform to turn a fixed-length collision-resistant compression function into a hash function that accepts arbitrary-length inputs. The input is "padded" so that its length is a multiple of 512 bits.

Compression Function - One-Way Function

one-way compression function is a function that transforms and mixes two fixed length inputs and produces a single fixed length output of the same size as one of the inputs. This can also be seen as that the compression function transforms one large fixed-length input into a shorter, fixed-length output. The transformation is "one-way", meaning that it is difficult given a particular output to compute inputs which compress to that output. One-way compression functions are not related to conventional data compression algorithms, which instead can be inverted exactly (lossless compression) or approximately (lossy compression) to the original data.

The mixing is done in such a way that full "avalanche effect" is achieved.

که این اثر بیانگر آن است که همه بیت های خروجی به تک تک بیت های ورودی وابسته است و هرگونه تغییر کوچکی در ورودی (تغییر حتی یک بیت) تاثیر قابل توجه و تصادفی در خروجی می گذارد. که این اثر در واقع همان ویژگی "نداشتن پیش زمینه ذهنی" از ورودی یا "مخفي سازی" را برای ما فراهم می کند.

Initialization Vector (IV), (Initial Hash Value)

"IV"، "مقدار اولیه"، یا "بردار اولیه"، یک مقدار 256-bit به عنوان مقدار اولیه تابع هش sha-256 میباشد، که به همراه اولین بلاک 512-bit از پیام، ورودی های اولین تابع فشرده سازی هستند. IV، شامل هشت کلمه 32-bit میشود که به عنوان Constant در هر بار محاسبه sha-256 به کار می رود. این مقادیر ثابت، در اسناد استاندارد مربوطه میتوان یافت. در شکل زیر مقادیر IV (در نمایش hex) آورده شده است.

$H_0^{(0)}$	=	6a09e667	$H_4^{(0)}$	=	510e527f
$H_1^{(0)}$	=	bb67ae85	$H_5^{(0)}$	=	9b05688c
$H_2^{(0)}$	=	3c6ef372	$H_6^{(0)}$	=	1f83d9ab
$H_3^{(0)}$	=	a54ff53a	$H_7^{(0)}$	=	5be0cd19

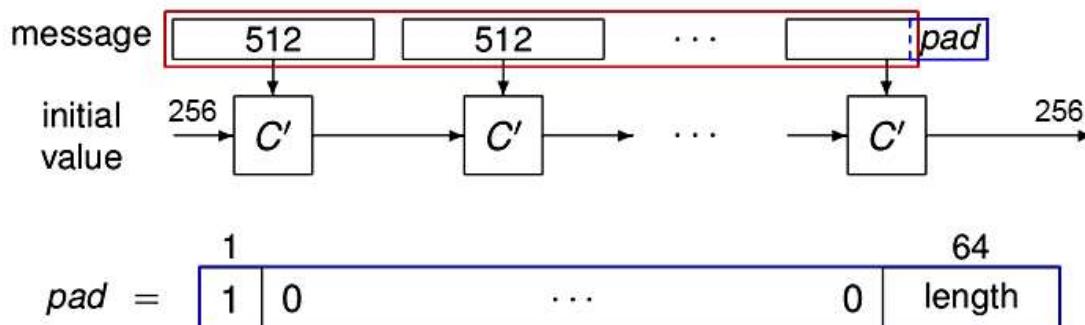
sha-256 مقادیر ثابت "IV" شامل هشت کمde (در نمایش hex) ثبت شده در استاندارد - Figure 5

These words were obtained by taking the first thirty-two bits of the fractional parts of the square roots of the first eight prime numbers.

Merkle–Damgard Strengthening & Padding

In order to make the construction secure, Merkle and Damgard proposed that messages be padded with a padding that encodes the length of the original message. This is called *length padding* or *Merkle–Damgard strengthening*.

The message isn't going to be, in general, necessarily exactly a multiple of the block size, so we're going to add some padding at the end of last block. And the padding is going to consist of, a 64-bit length field at the end of the padding, which is the length of the message in bits(binary). And then before that, a one bit, followed by some number of zero bits.



Note: maximum message length is $2^{64} - 1$ bits

نحوه sha-256 پلک آخر در ساختار Padding - Merkle-Damgard strengthening

First, the message is padded with a binary '1' then it is cut into blocks of 512 bits. If the length of the last block does not exceed 448 bits, as many zeros as necessary are appended to fill 448 bits and the binary length of the original message (before padding) is appended in the last 64 bits of the block to form a 512bit block. Else, the block is filled with zeros up to a length of 512 bits, and an extra block is appended filled with 448 zeros; again, the binary length of the original message is appended in the last 64 bits to form a complete 512-bit block. So, once you've padded the message such that, its length is exactly a multiple of the 512-bit block size. This form of padding is non-ambiguous and is an example of a valid Merkle-Damgard strengthening.

padding is always done, even if the message happens to be a multiple of the input block size. If padding is not always done, there is an easy hash collision, a message will have the same hash as that message with the pad appended.

2.2-HASH POINTERS AND DATA STRUCTURES

In this section, we'll discuss ways of using hash functions to build more complicated data structures that are used in distributed systems like Bitcoin. we're going to discuss **hash pointers** and their applications.

Hash pointer

A hash pointer is a data structure that is simply a pointer to where some information is stored together with a cryptographic hash of the value of that data at some fixed point in time. Whereas a regular pointer gives you a way to retrieve the information, a hash pointer also gives you a way to verify that the information hasn't changed.

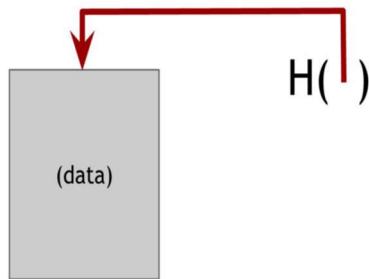


Figure 7 - A hash pointer.

Block chain

In Figure below, we built a linked list using hash pointers. We're going to call this data structure a **blockchain**. in a block chain, each block has data as well as a hash pointer to the previous block. So, each block not only tells us where the value of the previous block was, but it also contains a digest of that value that allows us to verify that the value hasn't changed. We store the head of the list, which is just a regular hash-pointer that points to the most recent data block.

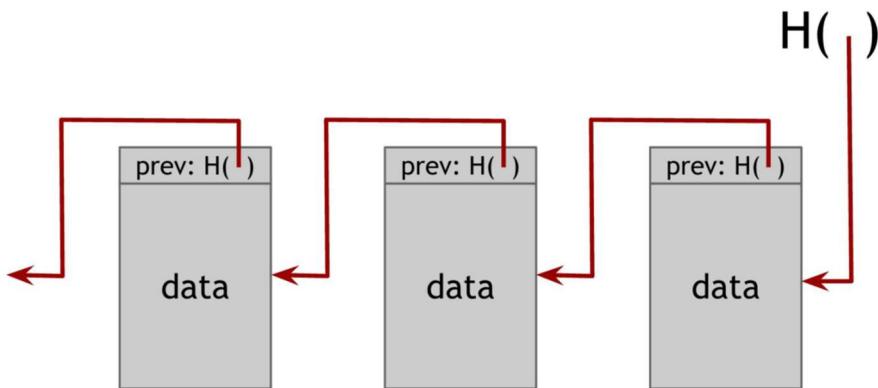


Figure 8 - A block chain is a linked list that is built with blocks of data and hash pointers

Application: tamper-evident log

A use case for a block chain is a tamper-evident log. That is, we want to build a log data structure that stores a bunch of data, and allows us to append data onto the end of the log. But if somebody alters data that is earlier in the log, we're going to detect it.

To understand why a block chain achieves this tamper-evident property, let's ask what happens if an adversary wants to tamper with data that's in the middle of the chain. the adversary changes the data of some block k. Since the data has been changed, the hash in block k + 1, which is a hash of the entire block k, is not going to match up. Remember that we are statistically guaranteed that the new hash will not match the altered content since the hash function is collision resistant. And so, we will detect

the inconsistency between the new data in block k and the hash pointer in block k + 1. the adversary can continue to try and cover up this change by changing the next block's hash as well. but this strategy will fail when he reaches the head of the list, and because he won't be able to tamper with that, as long as we store the hash pointer at the head of the list. Thus, by just remembering this single hash pointer, we've essentially remembered a tamper-evident hash of the entire list. and the adversary will be unable to change any block without being detected. So, we can build a block chain like this containing as many blocks as we want.

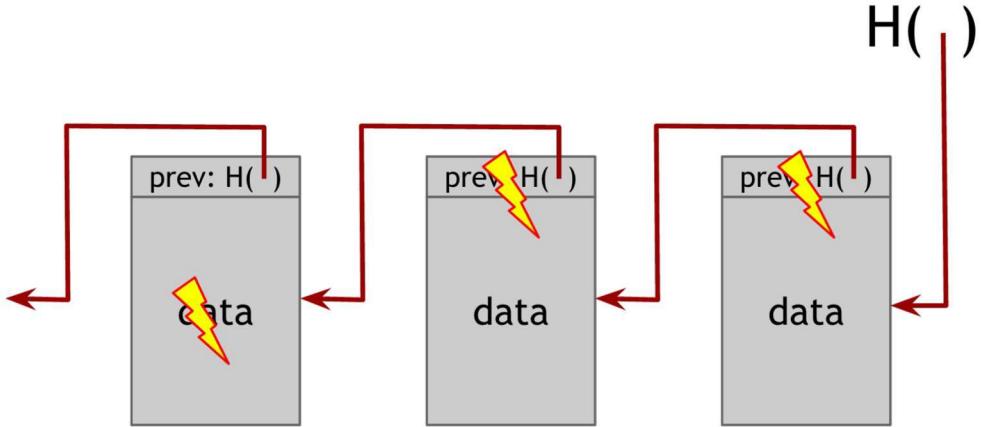


Figure 9 - Tamper-evident log. If an adversary modifies data anywhere in the block chain, it will result in the hash pointer in the following block being incorrect. If we store the head of the list, then even if the adversary modifies all of the pointers to be consistent with the modified data, the head pointer will be incorrect, and we will detect the tampering.

Merkle tree

Another useful data structure that we can build using hash pointers is a binary tree. A binary tree with hash pointers is known as a **Merkle tree**, after its inventor Ralph Merkle. Suppose we have a number of blocks containing data. These blocks comprise the leaves of our tree. We group these data blocks into pairs of two, and then for each pair, we build a data structure that has two hash pointers, one to each of these blocks. These data structures make the next level up of the tree. We in turn group these into groups of two, and for each pair, create a new data structure that contains the hash of each. We continue doing this until we reach a single block, the root of the tree.

As before, we remember just the hash pointer at the head of the tree. We now have the ability to follow paths through the hash pointers to any point in the list. This allows us to make sure that the data hasn't been tampered with because, just like we saw with the block chain, if an adversary tampers with some data block at the bottom of the tree, that will cause the hash pointer that's one level up to not match, and even if he continues to tamper with this block, the change will eventually propagate to the top of the tree where he won't be able to tamper with the hash pointer that we've stored. So again, any attempt to tamper with any piece of data will be detected by just remembering the hash pointer at the top.

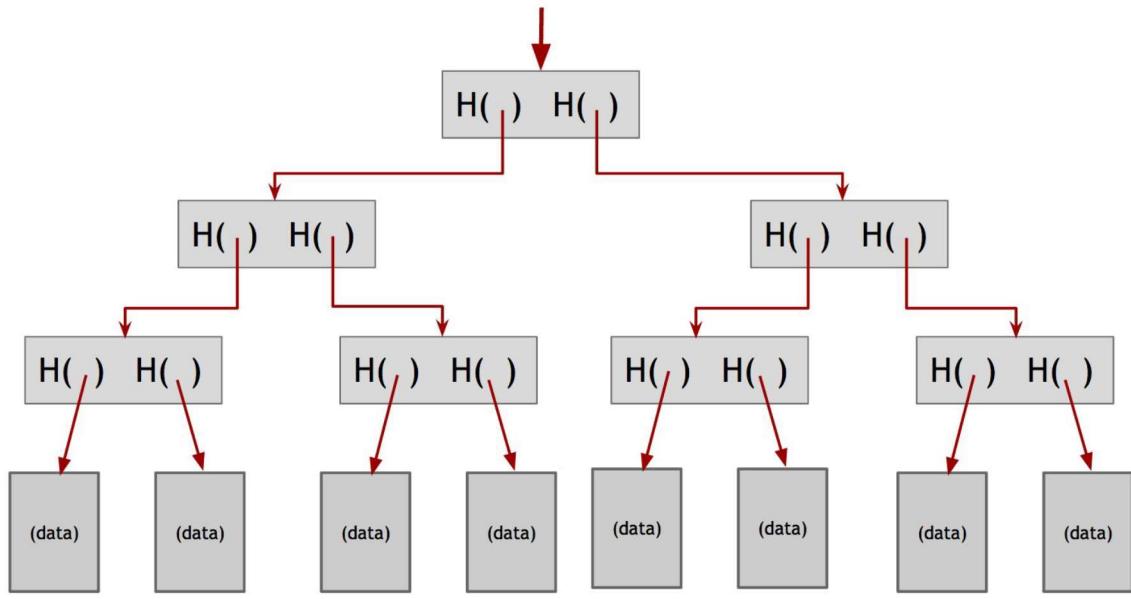


Figure 10 - **Merkle tree.** In a Merkle tree, data blocks are grouped in pairs and the hash of each of these blocks is stored in a parent node. The parent nodes are in turn grouped in pairs and their hashes stored one level up the tree. This continues all the way up the tree until we reach the root node.

Application1: proof of membership

Another nice feature of Merkle trees is that, unlike the block chain that we built before, it allows a concise proof of membership. Say that someone wants to prove that a certain data block is a member of the Merkle Tree. As usual, we remember just the root. Then they need to show us this data block, and the blocks on the path from the data block to the root. We can ignore the rest of the tree, as the blocks on this path are enough to allow us to verify the hashes all the way up to the root of the tree.

در حالی که در بلاک چین، چنین امکانی فراهم نبود. چراکه، جهت چک کردن وجود دیتای خاص در بلاک ها، نیاز به در اختیار داشتن: سایر دیتای همان بلاک، تمامی دیتای بلاک های بعد و هش پوینتر بلاک قبل است و همچنان تمام محاسبات مبایست کاملا تکرار شود.

See Figure below for a graphical depiction of how proof of membership in Merkle Tree works.

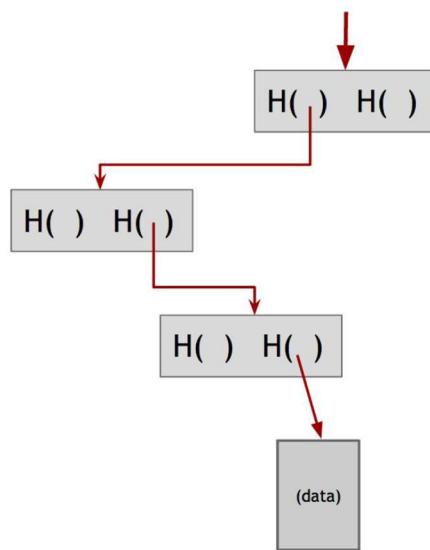


Figure 11 - **Proof of membership.** To prove that a data block is included in the tree, one only needs to show the blocks in the path from that data block to the root.

If there are n nodes in the tree, only about $\log(n)$ items need to be shown. And since each step just requires computing the hash of the child block, it takes about $\log(n)$ time for us to verify it. And so even if the Merkle tree contains a very large number of blocks, we can still prove membership in a relatively short time. Verification thus runs in time and space that's logarithmic in the number of nodes in the tree.

Application2: proof of non-membership - Variant: sorted Merkle tree

A **sorted Merkle tree** is just a Merkle tree where we take the blocks at the bottom, and we sort them using some ordering function. This can be alphabetical, lexicographical order, numerical order, or some other agreed upon ordering.

With a sorted Merkle tree, it becomes possible to verify non-membership in a logarithmic time and space. That is, we can prove that a particular block is not in the Merkle tree. And the way we do that is simply by showing a path to the item that's just before where the item in question would be and showing the path to the item that is just after where it would be. If these two items are consecutive in the tree, then this serves as a proof that the item in question is not included. For if it was included, it would need to be between the two items shown, but there is no space between them as they are consecutive.

2.3-DIGITAL SIGNATURES

In this section, we'll look at ***digital signatures***. This is the second cryptographic primitive, along with hash functions. A digital signature is an electronic analogue of a written signature. We desire two properties from digital signatures, Firstly, the digital signature can be used to provide assurance that the claimed signatory signed the information and only you can make your signature, but anyone who sees it can verify that it's valid. Secondly, a digital signature should be used to detect whether or not the information was modified after it was signed (i.e., to detect the integrity of the signed data). In other word, we want the signature to be tied to a particular document so that the signature cannot be used to indicate your agreement or endorsement of a different document. A properly implemented digital signature algorithm that meets the requirements of this Standard can provide these services.

Digital signature scheme

A digital signature scheme consists of the following three algorithms:

- **keys generation:** $(sk, pk) := \text{generatekeys(keysiz } e)$ The generatekeys method takes a key size and generates a key pair. The private key or secret key sk is kept privately and used to sign messages. Pk is the public verification key that you give to everybody. Anyone with this key can verify your signature.
- **Signature generation:** $\text{sig} := \text{sign}(sk , message)$ The sign method takes a message and a secret key, sk , as input and outputs a signature for $message$ under sk
- **Signature verification:** $\text{isValid} := \text{verify}(pk , message , sig)$ The verify method takes a message, a signature, and a public key as input. It returns a boolean value, $isValid$, that will be **true** if sig is a valid signature for $message$ under public key pk , and **false** otherwise.

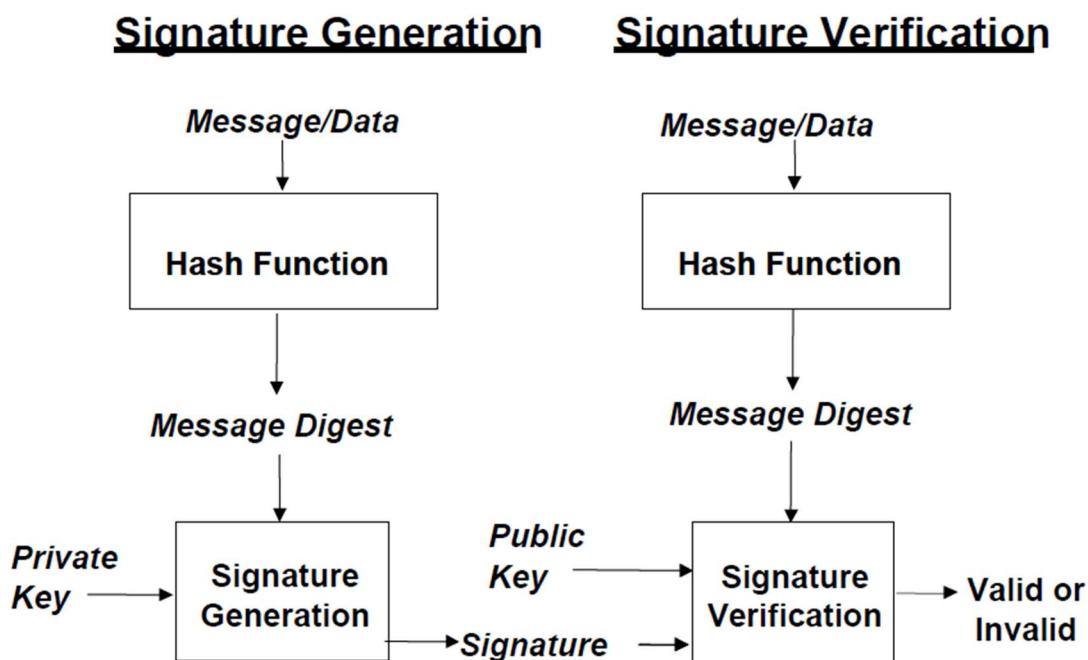


Figure 12 - Digital signature processes

A digital signature algorithm includes a signature generation process and a signature verification process. A signatory uses the generation process to generate a digital signature on data; a verifier uses the verification process to verify the authenticity of the signature. Each signatory has a public and

private key and is the owner of that key pair. the private key is used in the signature generation process. The key pair owner is the only entity that is authorized to use the private key to generate digital signatures. In order to prevent other entities from using the private key to generate fraudulent signatures, the private key must remain secret. The approved digital signature algorithms are designed to prevent an adversary who does not know the signatory's private key from generating the same signature as the signatory on a different message. The public key is used in the signature verification process. The public key need not be kept secret, anyone can verify a correctly signed message using the public key.

keys generation (Generatekeys) and **Signature generation (sign)** algorithms can be randomized algorithms. Indeed, generatekeys had better be randomized because it ought to be generating different keys for different people. **Signature verification**, on the other hand, will always be deterministic. For both the signature generation and verification processes, the message (i.e., the signed data) is converted to a fixed-length representation of the message by means of an approved hash function. Both the original message and the digital signature are made available to a verifier. A verifier also requires assurance that the key pair owner actually possesses the private key associated with the public key, and that the public key is a mathematically correct key.

Briefly, we require from these algorithms that the following two properties hold:

1. Valid signatures must verify: $\text{verify}(\text{pk}, \text{message}, \text{sign}(\text{sk}, \text{message})) == \text{true}$

If I sign a message with sk , my secret key, and someone later tries to validate that signature over that same message using my public key, pk , the signature must validate correctly. This property is a basic requirement for signatures to be useful at all.

2.Signatures are existentially unforgeable:

The second requirement is that it's computationally infeasible to forge signatures. That is, an adversary who knows your public key and gets to see your signatures on some other messages can't forge your signature on some message for which he has not seen your signature. Signature scheme is unforgeable if and only if, no matter what algorithm the adversary is using, his chance of successfully forging a message is extremely small, so small that we can assume it will never happen in practice.

ECDSA

Bitcoin uses a particular digital signature scheme that's called the Elliptic Curve Digital Signature Algorithm (ECDSA). ECDSA is a U.S. government standard, an update of the earlier DSA algorithm adapted to use elliptic curves. These algorithms have received considerable cryptographic analysis over the years and are generally believed to be secure.

More specifically, Bitcoin uses ECDSA over the standard elliptic curve "secp256k1" which is estimated to provide 128 bits of security (that is, it is as difficult to break this algorithm as performing 2^{128} symmetric-key cryptographic operations such as invoking a hash function).

It might be useful to have an idea of the sizes of various quantities in ECDSA:

Private key: 256 bits

Public key, uncompressed: 512 bits

Public key, compressed: 257 bits

Message to be signed: 256 bits

Signature: 512 bits

Figure 13 - sizes of various quantities in ECDSA

Note that while ECDSA can technically only sign messages 256 bits long, this is not a problem: messages are always hashed before being signed, so effectively any size message can be efficiently signed. Also, you can sign a hash pointer. If you sign a hash pointer, then the signature covers, or protects, the whole structure and everything the chain of hash pointers points to. For example, if you were to sign the hash pointer that was at the end of a block chain, the result is that you would effectively be digitally signing the entire block chain.

Public Keys as Identities

The idea is to take a public key and equate that to an identity of a person or an actor in a system. If you see a message with a signature that verifies correctly under a public key, pk , then you can think of this as pk is saying the message. via the **generateKeys** operation in our digital signature scheme. pk is the new public identity that you can use, and sk is the corresponding secret key that only you know and lets you speak for on behalf of the identity pk . In practice, you may use the hash of pk as your identity since public keys are large. If you do that, then in order to verify that a message comes from your identity, one will have to check (1) that pk indeed hashes to your identity, and (2) the message verifies under public key pk . Moreover, by default, your public key pk will basically look random, and nobody will be able to uncover your real-world identity by examining pk . But of course, once you start making statements using this identity, these statements may leak information that allows one to connect pk to your real-world identity.

Application: Decentralized identity management.

This brings us to the idea of decentralized identity management. Rather than having a central authority that you have to go to in order to register as a user in a system, you can register as a user all by yourself. You don't need to be issued a username nor do you need to inform someone that you're going to be using a particular name. If you want to be somewhat anonymous for a while, you can make a new identity, use it just for a little while, and then throw it away. All of these things are possible with decentralized identity management, and this is the way Bitcoin, in fact, does identity. These identities are called **addresses**, in Bitcoin jargon. And that's really just a hash of a public key.

Pseudo-anonymity: in Bitcoin you don't need to explicitly register or reveal your real-world identity, but the pattern of your behavior might itself be identifying, this feature called Pseudo-anonymity.

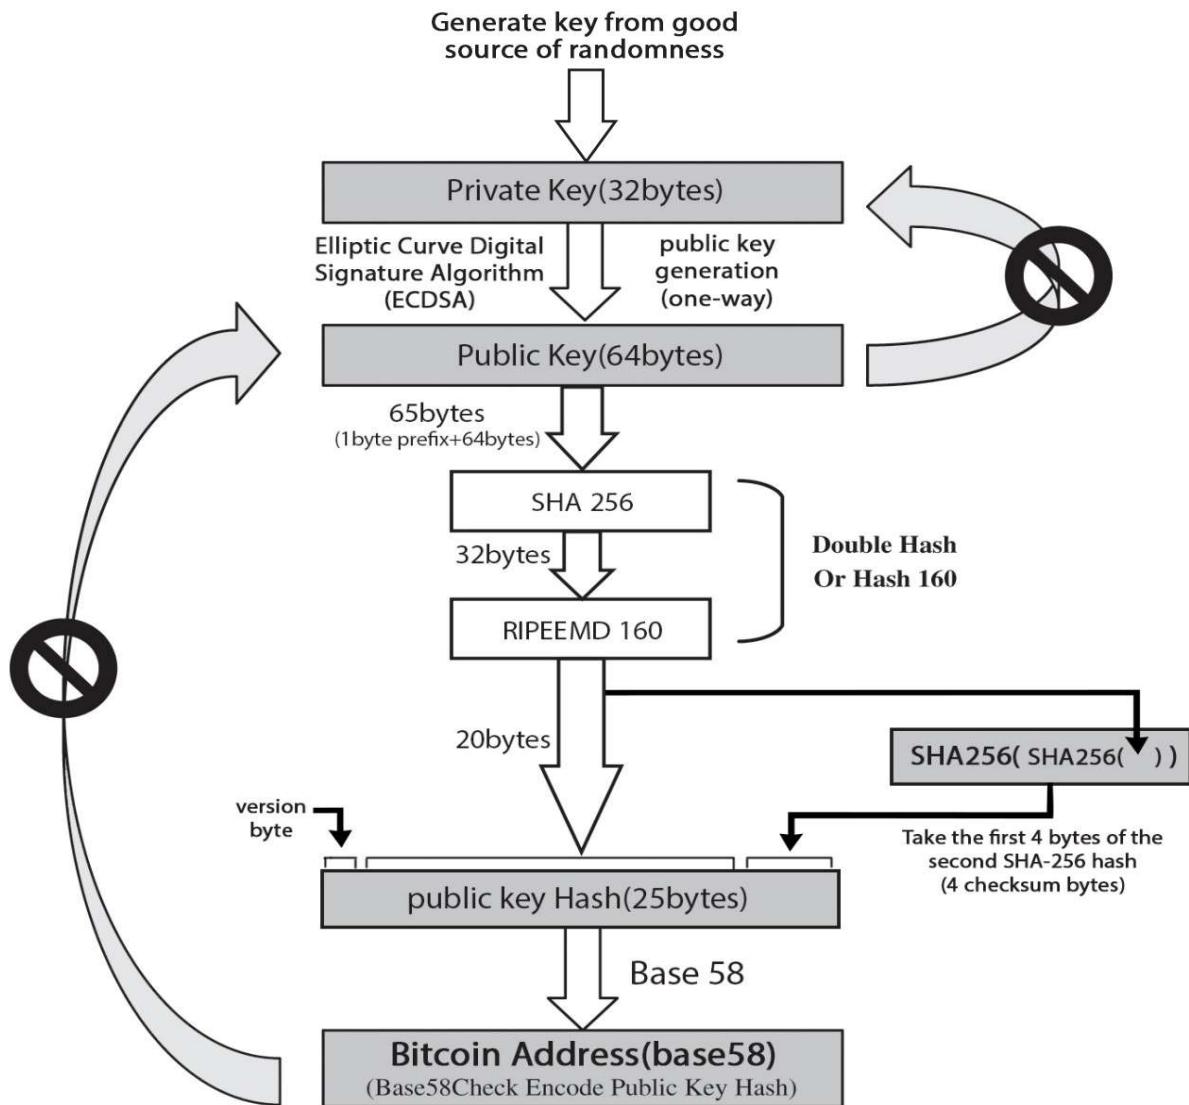


Figure 14 - Conversion from private key to Bitcoin Address

3-BITCOIN MINING

Bitcoin, at a simple glance, is three things. First it is a protocol (or set of rules) that defines how the network should operate. Second it is a software project that implements that protocol. Third it is a network of computers and devices running software that uses the protocol to create and manage the Bitcoin currency. There's no central authority for Bitcoins, similar to a central bank which controls currencies. Instead, programmers solve complex puzzles to endorse Bitcoin transactions and get Bitcoins as a reward. This activity is called Bitcoin mining. Mining is defined in the protocol, implemented in software, and is an essential function in managing the Bitcoin network. Miners validate every transaction, they build and store all the blocks, and they reach a consensus on which blocks to include in the block chain.

3.1-The task of Bitcoin miners

To be a Bitcoin miner, you have to join the Bitcoin network and connect to other nodes. Once you're connected, there are six tasks to perform:

1. ***Listen for transactions.*** First, you listen for transactions on the network and validate them by checking that signatures are correct and that the outputs being spent haven't been spent before.
2. ***Maintain block chain and listen for new blocks.*** You must maintain the block chain. You start by asking other nodes to give you all of the historical blocks that are already part of the block chain before you joined the network. You then listen for new blocks that are being broadcast to the network. You must validate each block that you receive by validating each transaction in the block and checking that the block contains a valid nonce. We'll return to the details of nonce checking later in this section.
3. ***Assemble a candidate block.*** Once you have an up-to-date copy of the block chain, you can begin building your own blocks. To do this, you group transactions that you heard about into a new block that extends the latest block you know about. You must make sure that each transaction included in your block is valid.
4. ***Find a nonce that makes your block valid.*** This step requires the most work and it's where all the real difficulty happens for miners.
5. ***Hope your block is accepted.*** Even if you find a block, there's no guarantee that your block will become part of the consensus chain. There's bit of luck here; you have to hope that other miners accept your block and start mining on top of it, instead of some competitor's block.
6. ***Profit.*** If all other miners do accept your block, then you profit! At the time of this writing in early 2022, the block reward is 6.25 bitcoins which is currently worth over \$312,500. In addition, if any of the transactions in the block contained transaction fees, the miner collects those too. Mathematically, transaction fees are the difference between the amount of bitcoin sent and the amount received. Conceptually, transaction fees are a reflection of the speed with which a user wants their transaction validated on the blockchain. Transaction fees are based on the data volume of a transaction and the congestion of the network. through the Segregated Witness (SegWit) protocol upgrade in 2017, now A block can contain a maximum theoretical limit of 4 MB of data and a more realistic limit of 2 MB. The default limit before the upgrade was 1MB. so, there is a limit to how many transactions can be processed in one block. A larger transaction will take up more block data. Thus, larger transactions typically pay fees on a per-byte basis. If you wish to have your transaction confirmed immediately, your optimal fee rate may vary significantly.

We can classify the steps that a miner must take into two categories.

1. **Validating transactions and blocks.** Some tasks help the Bitcoin network and are fundamental to its existence. These tasks are the reason that the Bitcoin protocol requires miners in the first place
2. **The race to find blocks and profit.** Some tasks aren't necessary for the Bitcoin network itself but are intended to incentivize miners to perform the essential steps. Of course, both of these are necessary for Bitcoin to function as a currency, since miners need an incentive to perform the critical steps.

Finding a valid block

In previous sections, we saw that there are two main hash-based structures. There's the block chain where each block header points to the previous block header in the chain, and then within each block there's a Merkle tree of all of the transactions included in that block. The first thing that you do as a miner is to compile a set of valid transactions that you have from your pending transaction pool into a Merkle tree. Of course, you may choose how many transactions to include up to the limit on the total size of the block. You then create a block with a header that points to the previous block. In the block header, there's a 32-bit nonce field, and you keep trying different nonces looking for one that causes the block's hash to begin with the required number of zeros. A miner may begin with a nonce of 0 and successively increment it by one in search of a nonce that makes the block valid.

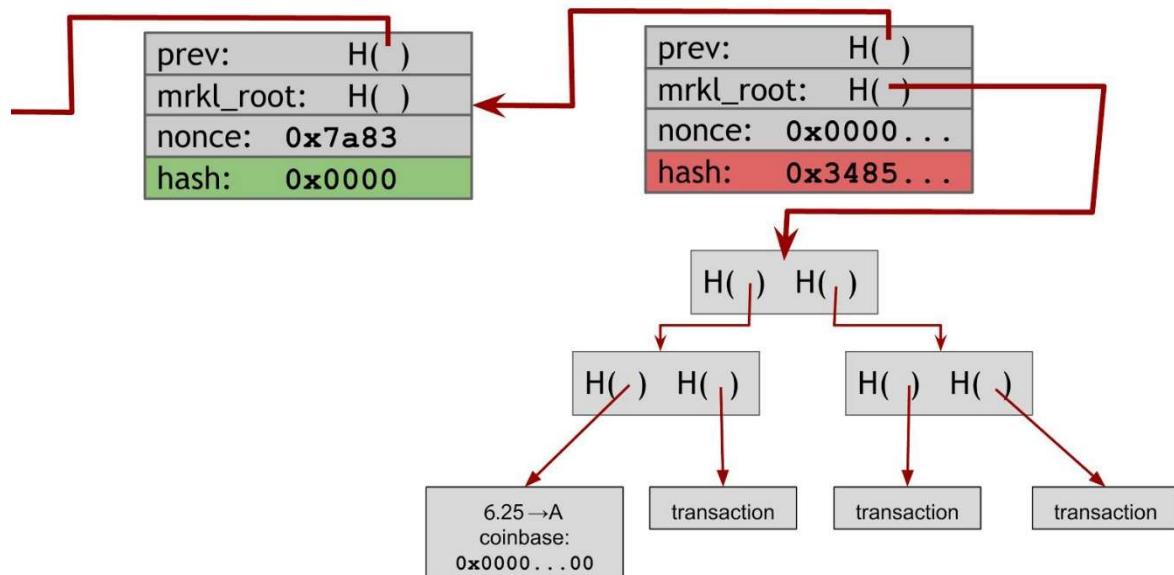


Figure 15 - Finding a valid block. In this example, the miner tries a nonce of all 0s. It does not produce a valid hash output, so the miner would then proceed to try a different nonce.

In most cases you'll try every single possible 32-bit value for the nonce and none of them will produce a valid hash. At this point you're going to have to make further changes. Notice in Figure above that there's an additional nonce in the Coinbase transaction that you can change as well. After you've exhausted all possible nonces for the block header, you'll change the extra nonce in the Coinbase transaction, for example, by incrementing it by one, and then you'll start searching nonces in the block header once again.

When you change the nonce parameter in the Coinbase transaction, the entire Merkle tree of transactions has to change (See Figure below). Since the change of the Coinbase nonce will propagate all the way up the tree, changing the extra nonce in the Coinbase transaction is much more expensive operation than changing the nonce in the block header. For this reason, miners spend most of their time changing the nonce in the block header and only change the Coinbase nonce when they have exhausted all of the 2^{32} possible nonces in the block header without finding a valid block.

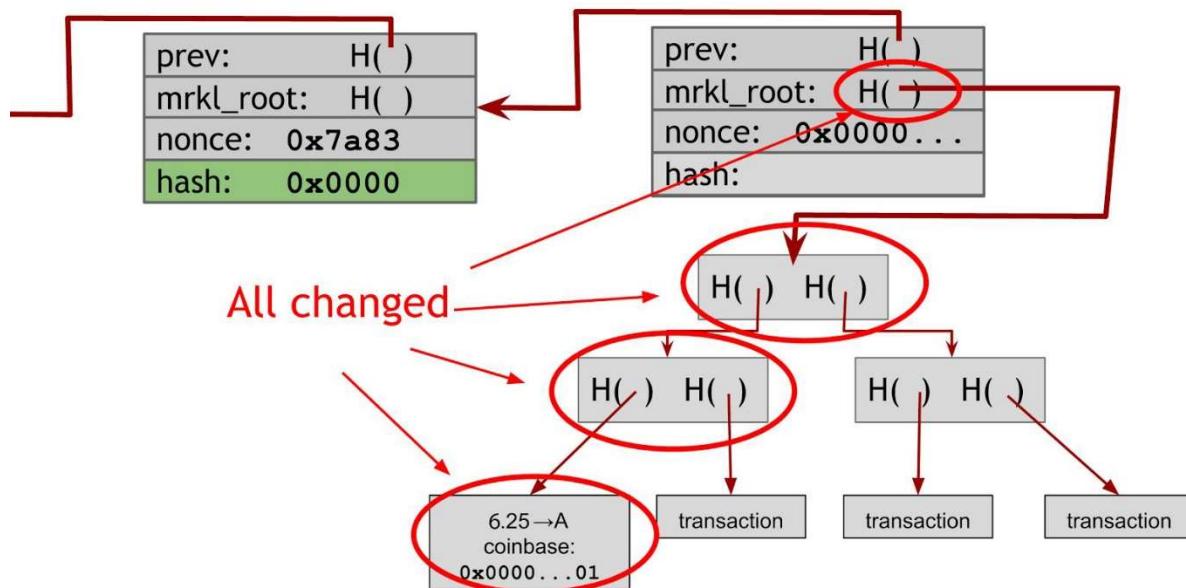


Figure 15 - Changing a nonce in the Coinbase transaction propagates all the way up the Merkle tree.

The vast, vast majority of nonces that you try aren't going to work, but if you stay at it long enough, you'll eventually find the right combination of the extra nonce in the Coinbase transaction and the nonce in the block header that produce a block with a hash under the target. When you find this, you want to announce it as quickly as you can and hope that you can profit from it.

Is everyone solving the same puzzle?

If every miner just increments the nonces as described, aren't all miners solving the exact same puzzle? Won't the fastest miner always win? The answer is no! Firstly, it's unlikely that miners will be working on the exact same block as each miner will likely include a somewhat different set of transactions and in a different order. But more importantly, even if two different miners were working on a block with identical transactions, the blocks would still differ. Recall that in the Coinbase transaction, miners specify their own address as the owner of the newly minted coins. This address by itself will cause changes which propagate up to the root of the Merkle tree, ensuring that no two miners are working on exactly the same puzzle unless they share a public key. This would only happen if the two miners are part of the same mining pool (which we'll discuss shortly), in which case they'll communicate to ensure they include a distinct nonce in the Coinbase transaction to avoid duplicating work.

Difficulty

The Bitcoin mining difficulty defines how hard it is to mine a new block. In other word, Difficulty is a measure of how difficult it is to find a hash below a given target. The Bitcoin network has a global block difficulty. Valid blocks must have a hash below this target.

There are two values we need to consider:

Target ("Bits"): The Bitcoin difficulty target is a 256-bit number. It is stored as an encoded packed representation (called "Bits") for its actual hexadecimal target. The target states how high a block hash can be in order to fulfill the mining condition. The higher the target the easier it is to find a value below it and the lower the difficulty.

Difficulty: Is a numerical expression of how difficult it is to find a suitable hash compared to the easiest difficulty of 1. That means, Floating point representation of difficulty shows how much current target is harder than the one used in the genesis block.

Figure below shows summary information of first block mined in 2022:

Summary

Block Hash	000000000000000000000000001d075e7b9eba4b03ce577325723a819767a9e4a85aa84		
Height	◀ 716,572 ▶	Block Reward	6.25000000 BTC
Time	2022-01-01 00:02:12	Fee Reward	0.05391132 BTC
Block Size	1,557,044 Bytes	Tx Count	2,694
Relayed By	 Foundry USA	Tx Volume	2,500.45384565 BTC
Difficulty	24.27 T		

Block Header

Merkle Root	b66e82f405d605850a64271dc5e16e7d0d511b57803d1fef99ccab6ec687e8
Version	0x20600000
Nonce	0x53811281
Bits	0x170b98ab

Figure 16 - summary information of first block mined in 2022.

For this block, as of January 2022, figure below shows how to decode the target "bits" and determine the mining difficulty.

Bit as hexadecimal number:

0x170b98ab

Coefficient: 0x17

Index: 0x0b98ab

Target as hexadecimal with padding:

Target as hexadecimal:

Target as decimal number:

1110710552837102268873518195482888052995095958078357504

Difficulty as decimal number:

24272331996979 R 1077885226062650189330663434140080336373623713163116544

$$\log_2 \left(\frac{2^{256}}{1110710552837102268873518195482888052995095958078357504} \right) = 76.4643999715$$

Figure 17 - decode the target "bits" and determining proses for the mining difficulty and target

So, the hash of any valid block has to be below this value. In other words, one in about 2^{76} nonces (Based on calculation shows in figure above) that you try will work, which is a really huge number. One approximation is that it's 1680-times greater than the human population of earth squared. So, if every person on earth was themselves their own planet earth with 7.8 billion people on it, the total number of people multiply by 1680, would be close to 2^{76} .

Determining the difficulty

Difficulty gets adjusted according to the hash rate in order to maintain a constant block finding time and to dispense the block reward not too fast or too slow. The mining difficulty changes every 2016 blocks, which are found about once every 2 weeks. It is adjusted based on how efficient the miners were over the period of the previous 2016 blocks according to this formula:

`next_difficulty = (previous_difficulty * 2016 * 10 minutes) / (time to mine last 2016 blocks)`

Note that $2016 * 10$ minutes is exactly two weeks, so 2016 blocks would take two weeks to mine 2016 blocks if a block were created exactly every 10 minutes. So, the effect of this formula is to scale the difficulty to maintain the property that blocks should be found by the network on average about once every ten minutes. There's nothing special about 2 weeks, but it's a good trade-off. If the period were much shorter, the difficulty might fluctuate due to random variations in the number of blocks found in each period. If the period were much higher, the network's hash power might get too far out of balance with the difficulty.

You can see in Figure below that over time the mining difficulty keeps increasing. It's not necessarily a steady linear increase or an exponential increase, but it depends on activity in the market. Mining difficulty is affected by factors like how many new miners are joining, which in turn may be affected by the current exchange rate of Bitcoin. Generally, as more miners come online and mining hardware gets more efficient, blocks are found faster and the difficulty is increased so that it always takes about ten minutes to find a block.

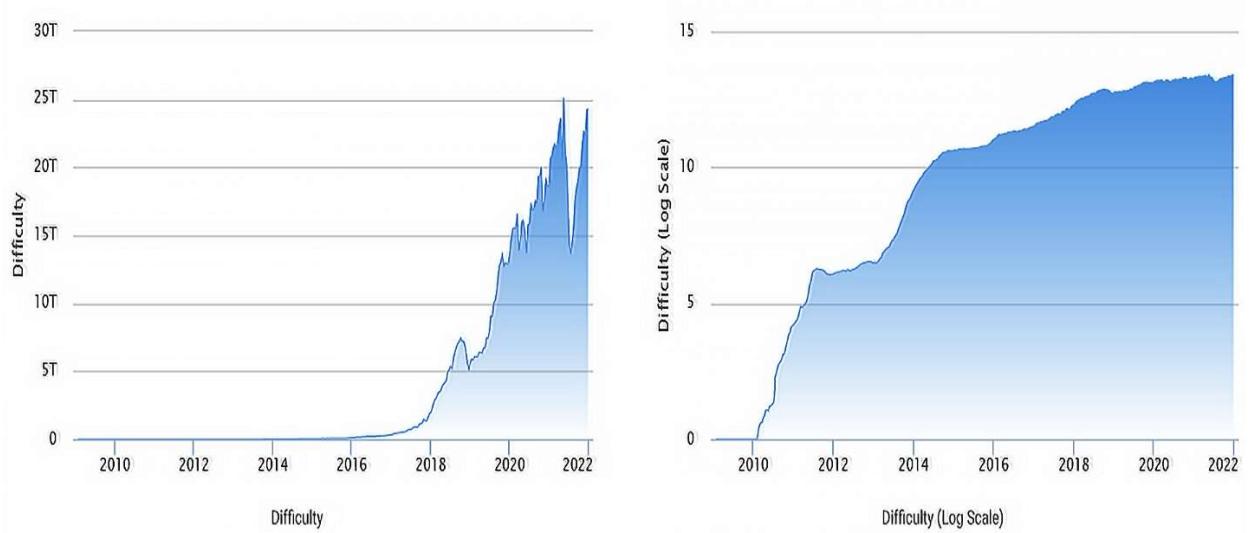


Figure 18 - Difficulty over time - left: over actual scale time - Right: over Log scale time

Each Bitcoin miner independently computes the difficulty and will only accept blocks that meet the difficulty that they computed. Miners who are on different branches might not compute the same difficulty value, but any two miners mining on top of the same block will agree on what the difficulty should be. This allows consensus to be reached.

3.2-Mining Hardware

Notice that the computation that miners have to do is very difficult. In this section, we'll discuss why it is so computationally difficult and take a look at the hardware that miners use to perform this computation.

The core of the difficult computation miners are working on is the SHA-256 hash function. We discussed hash functions and properties of those in previous sections. SHA-256 is a general-purpose cryptographic hash function that's part of a bigger family of functions that was standardized in 2001 (SHA stands for Secure Hash Algorithm). SHA-256 was a reasonable choice as this was strongest cryptographic hash function available at the time when Bitcoin was designed. It is possible that it will become less secure over the lifetime of Bitcoin, but for now it remains secure.⁵

A closer look at SHA-256

Figure below shows more detail about what actually goes on in a SHA-256 computation. SHA-256 maintains 256 bits of state(variable). The state(variable) is split into eight 32-bit words(a-b-c-d-e-f-g-h) which makes it highly optimized for 32-bit hardware. for first 512-bit block of message these eight 32-bit words are Initialization Vectors (IV). and for subsequent blocks of message, 256 bits of the state(variable) are 256-bit hash of previous block.

At a very simple glance, for each 512-bit block of message (or for each compression function), message prepare and schedule of sixty-four 32-bit words. in each round of 64 round, a number of words in the state are taken (some with bitwise tweaks related to words of message applied) and added together mod 32. The entire state is then shifted⁶ over with the result of the addition becoming the new left most word of the state (the result value placed as new eight 32-bit words of state).

this process repeated 64-time for each block of message. And the latest result of eight 32-bit words of state, forms the hash of the block.

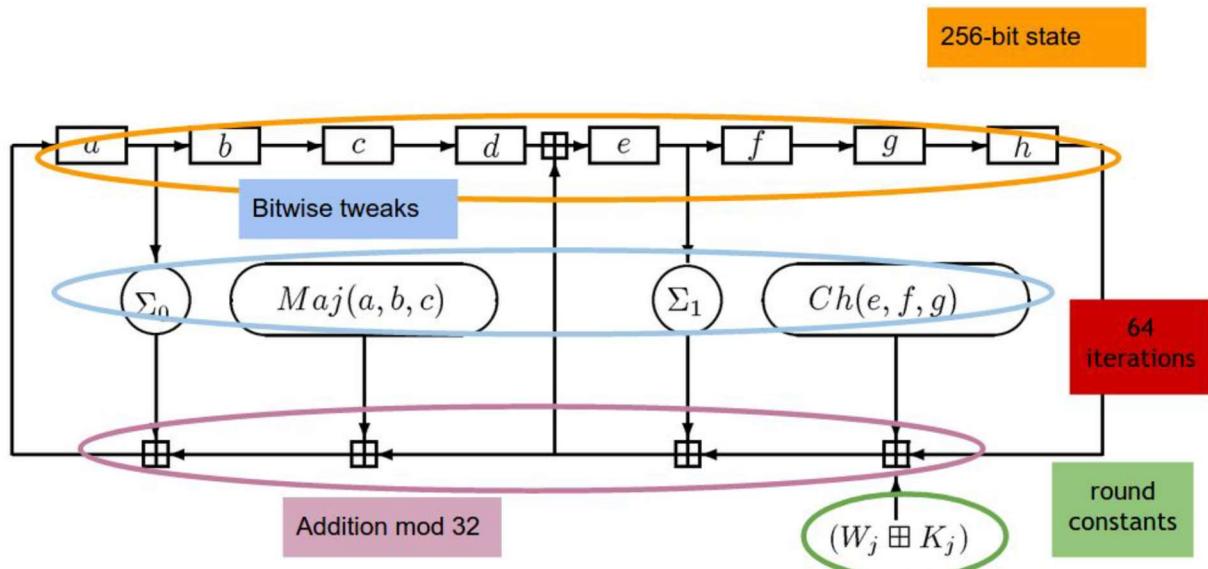


Figure 19 - The structure of SHA-256. This is one round of the compression function.

⁵ Although the SHA-2 family, including SHA-256, are still considered to be cryptographically secure, the next generation SHA-3 family has now been picked by a contest. SHA-3 is in the final stages of standardization today, but it wasn't available at the time Bitcoin was designed.

⁶ This part of design is inspired by bitwise Linear Feedback Shift Registers (LFSR).

Figure above shows just one round of the SHA-256 compression function. A complete computation of SHA-256 does this for 64 iterations. During each round, there are slightly different constants applied so that no iteration is exactly the same.

The task for miners is to compute this function as quickly as possible. Remember that miners are racing each other so the faster they do this, the more they earn. Bitcoin actually requires SHA-256 to be applied twice to a block in order to get the hash that is used by the nodes. The reasons for the double computation are not fully specified, but at this point, it's just something that miners have to deal with.

CPU mining

The first generation of mining was all done on general purpose computers that is general purpose central processing units (CPUs). In fact, CPU mining was running a pseudocode like the code shown in Figure below. That is, miners searched over nonces in a linear procedure, computed SHA 256 in software and checked if the result was a valid block.

```
TARGET = (65535 << 208) / DIFFICULTY;
coinbase_nonce = 0;
while (1) {
    header = makeBlockHeader(transactions, coinbase_nonce);
    for (header_nonce = 0; header_nonce < (1 << 32); header_nonce++){
        if (SHA256(SHA256(makeBlock(header, header_nonce))) <
TARGET)
            break; //block found!
    }
    coinbase_nonce++;
}
```

Figure 20 - CPU mining pseudocode.⁷

How fast will this run on a general-purpose computer? On a good desktop PC, you might expect to compute about on scale by 50 million hashes per second (MH/s).

If you're mining on a general-purpose PC today, CPU mining is no longer profitable with the current difficulty. For the last few years, anyone trying to mine on a CPU probably doesn't understand how Bitcoin works and was probably pretty disappointed that they never made any money doing it.

GPU mining

The second generation began when people started to get frustrated with how slow their CPUs were and instead used their graphics card, or graphics processing unit (GPU).

Bitcoin mining can be parallelized easily because you can try computing multiple hashes at the same time with different nonces. In 2010, a language called OpenCL was released. OpenCL is a general-purpose language to do things other than graphics on a GPU. It's a high level-language and over time people have used it to run many types of computation more quickly on graphics cards. This paved the way for Bitcoin mining on GPUs.

Mining with graphics cards had several attractive properties at the time. For one thing, they're easily available and easy for amateurs to set up. They're the most accessible high-end hardware that's available to the general public. They also have some properties that make them specifically good for Bitcoin mining. They're designed for parallelism so they have many Arithmetic Logic Units (ALUs) that

⁷ notice in the code that as mentioned, SHA-256 is applied twice.

can be used for simultaneous SHA-256 computations. Some GPUs also have specific instructions to do bitwise operations that are quite useful for SHA-256.

Most graphics cards can also be **overclocked**, meaning you can run them faster than they're actually designed for if you want to take on the risk that they might overheat or malfunction. With Bitcoin mining, it might be profitable to run the chip much faster than it was designed for even if you induce a few errors by doing so. For example, consider you can run your graphics card 50 percent faster but doing so will cause errors in the SHA-256 computation to 30 percent of the time. In the above example, the throughput is 1.5x compared to not overclocking, whereas the success rate is 0.7x. The product is 1.05, which means overclocking increases your expected profits by 5%.

Finally, you can drive many graphics cards from one motherboard and CPU. So, you can take your computer, which will be running your actual Bitcoin node which gathers transactions from the network and assembles blocks, and attach multiple graphics cards to it to try to find the right nonces to make the SHA-256 of the block valid. Figure below shows a setup to drive many, many GPUs from a single CPU. This was still in the early days of Bitcoin when miners were still mostly hobbyists without much experience running servers, but they came up with some quite ingenious designs for how to pack many graphics cards into a small place and keep them cool enough to operate.



Figure 21 - A home-built rack of GPUs used for Bitcoin mining. You can also see the fans that they used to build a primitive cooling system.

On a good graphics card with aggressive tuning, you might get on scale by 500 MH/s, or 500 million hashes per second.

Disadvantages of GPU mining

GPU mining has some disadvantages. GPUs have a lot of hardware built into them for doing video processing that can't be utilized for mining. Specifically, they have a large number of floating-point units that aren't used at all in SHA-256. GPUs also don't have the greatest cooling characteristics when you put a lot of them next to one another. They're not designed to run side by side as they are in the picture; they're designed to be in a single box doing graphics for one computer.

GPUs can also draw a fairly large amount of power, so a lot of electricity is used relative to a computer. Another disadvantage initially was that you had to either build your own board or buy expensive boards to house multiple graphics cards.

FPGA mining

Around 2011 some miners started switching from GPUs to FPGAs, or Field Programmable Gate Arrays, after the first implementation of Bitcoin mining came out in Verilog, a hardware design language that's used to program FPGAs. The general rationale behind FPGAs is to try to get close as possible to the performance of custom hardware while also allowing the owner of the card to customize it or reconfigure it "in the field." By contrast, custom hardware chips are designed in a factory and do the same thing forever.

FPGAs offer better performance than graphics cards, particularly on "bit fiddling" operations which are trivial to specify on an FPGA. Cooling is also easier with FPGAs and, unlike GPUs, you can theoretically use nearly all of the transistors on the card for mining. Like with GPUs, you can pack many FPGAs together and drive them from one central unit, which is exactly what people began to do (see Figure below). Overall, it was possible to build a big array of FPGAs more neatly and cleanly than you could with graphics cards.



Figure 22 - A home-built rack of FPGAs

Using an FPGA with a careful implementation, you might get up to several GH/s billion hashes per second. This is certainly a large performance gain over CPUs and GPUs.

Despite the performance gain, the days of FPGA mining were quite limited. Firstly, they were being driven harder for Bitcoin mining — by being on all the time and overclocked — than consumer grade FPGAs were really designed for. Because of this, many people saw errors and malfunctions in their FPGAs as they were mining. It also turned out to be difficult to optimize the 32-bit addition step which is critical in doing SHA-256. Secondly, FPGAs are also less accessible—you can't buy them at most stores and there are fewer people who know how to program and set up an FPGA than a GPU.

Most importantly though, even though FPGAs improved performance the cost-per-performance was only marginally improved over GPUs. This made FPGA mining was a rather short-lived phenomenon. Whereas GPU mining dominated for about a year.

ASIC⁸ mining

Mining today is dominated by Bitcoin ASICs, or **application-specific integrated circuits**. These are chips that were designed, built, and optimized for the sole purpose of mining Bitcoins. There are a few big vendors that sell these to consumers with a good deal of variety: you can choose between slightly bigger and more expensive models, more compact models, as well as models with varying performance and energy consumption claims.

The performance of today's miners is about above of 1,000GH/s one trillion hashes per second. This is certainly a large performance gain over CPUs and GPUs and FPGAs.

Designing ASICs requires considerable expertise and their lead-time is also quite long. Nevertheless, Bitcoin ASICs were designed and produced surprisingly quickly. In fact, analysts have said that this may be the fastest turnaround time in the history of integrated circuits from specifying a problem and to have a working chip in people's hands. Partially as a result of this, the first few generations of Bitcoin ASICs were quite buggy and most of them didn't quite deliver the promised performance numbers. Bitcoin ASICs have since matured and there are now fairly reliable ASICs available.

Up until 2014, the lifetime of ASICs was quite short due to the rapidly increasing network hash rate, with most boards in the early ASIC era growing obsolete in about six months. Within this time, the bulk of the profits are made up front. Often, miners will make half of the expected profits for the lifetime of the ASIC during just the first six weeks. This meant, even shipping speed can become a crucial factor in making a profit. As the growth rate of Bitcoin's hash power has stabilized, mining equipment has a longer life time.



Figure 23 - A commercial ASIC bitcoin miner

⁸ ASIC stands for application specific integrated circuit, which is a specialized silicon chip that performs just one task. In the digital currency space, these chips process SHA-256 in order to mine bitcoin and validate transactions.

4-FPGA

4.1-Introduction

Field programmable gate array (FPGA) is a type of integrated circuit (IC) that can be programmed for different algorithms after fabrication. FPGA first introduced in 1985. Modern FPGA devices consist of up to four million logic cells that can be configured to implement a variety of software algorithms. an FPGA provides significant cost advantages in comparison to an IC development effort and offers the same level of performance in most cases. FPGA is developed version of PLD⁹. In PLD'S programmed after manufacturing in field, it has limited programmability. Field programmable gate is capable of implementing any digital circuit. This provides developer of creating wide array of logical structure minimum low cost.



Figure 24 - A sample FPGA - Xilinx spartan series

Requirement of FPGAs

By the early 1980's large scale integrated circuits (LSI) formed the back bone of most of the logic circuits in major systems. Microprocessors, bus/IO controllers, system timers etc. were implemented using integrated circuit fabrication technology. custom "glue logic"¹⁰ or interconnects were still required to help connect the large integrated circuits in order to:

Generate global control signals (for resets etc.)

Data signals from one subsystem to another sub system.

First attempt to this problem was developing custom ICs, but it was not a solution for the problem. custom ICs have their own disadvantages. They are relatively very expensive to develop, and delay introduced for product to market (time to market) because of increased design time. There are two kinds of costs involved in development of custom ICs

1. Cost of development and design
2. Cost of manufacture

⁹ PLDs are array-oriented devices that typically have an AND-OR structure with wide-input AND gates feeding a narrower OR gate. The difference between FPGA and PLD is that **FPGA incorporates logic blocks instead of fixed AND-OR gates**. FPGAs are designed for having higher gate count whereas, PLDs are used for lesser gate counts.

¹⁰ Glue logic is a special form of digital circuitry that allows different types of logic chips or circuits to work together by acting as an interface between them.

FPGAs were introduced as an alternative to custom ICs for implementing entire system on one chip and to provide flexibility of reprogramability to the user. One of the major advantages of FPGA over custom ICs is circuit implementation time is very less because physical layout, masking etc. is absent in this process. Circuit implementation is done with help of the advanced CAD¹¹ tools.

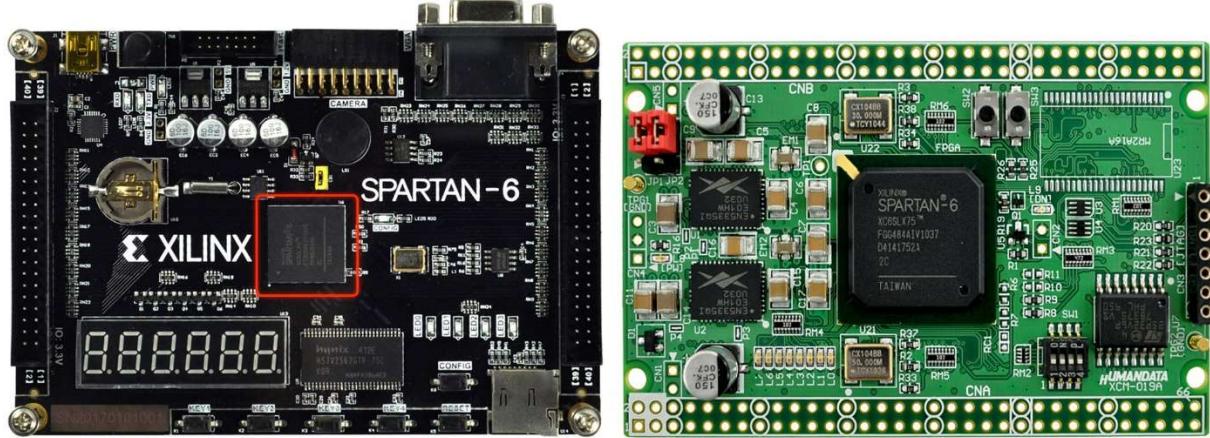


Figure 25 – FPGA Development Boards -Xilinx Spartan 6

¹¹ Computer-aided design (CAD) is the use of computers to aid in the creation, modification, analysis, or optimization of a design.

4.2-Architecture and Structure of FPGA

FPGA is a device that contains a matrix of reconfigurable gate array logic circuitry. When a FPGA is configured, the internal circuitry is connected in a way that creates a hardware implementation of the software application. Unlike processors, FPGAs use dedicated hardware for processing logic and do not have an operating system. FPGAs are truly parallel in nature so different processing operations do not have to compete for the same resources. As a result, the performance of one part of the application is not affected when additional processing is added. Unlike ASICs and hard-wired printed circuit board (PCB) designs which have fixed hardware resources, FPGA-based systems can literally rewire their internal circuitry to allow reconfiguration after the control system is deployed to the field.

A single FPGA can replace thousands of discrete components by incorporating millions of logic gates in a single integrated circuit (IC) chip. The internal resources of an FPGA chip consist of a matrix of configurable logic blocks (CLBs) surrounded by a periphery of I/O blocks. Signals are routed within the FPGA matrix by programmable interconnect switches and wire routes.

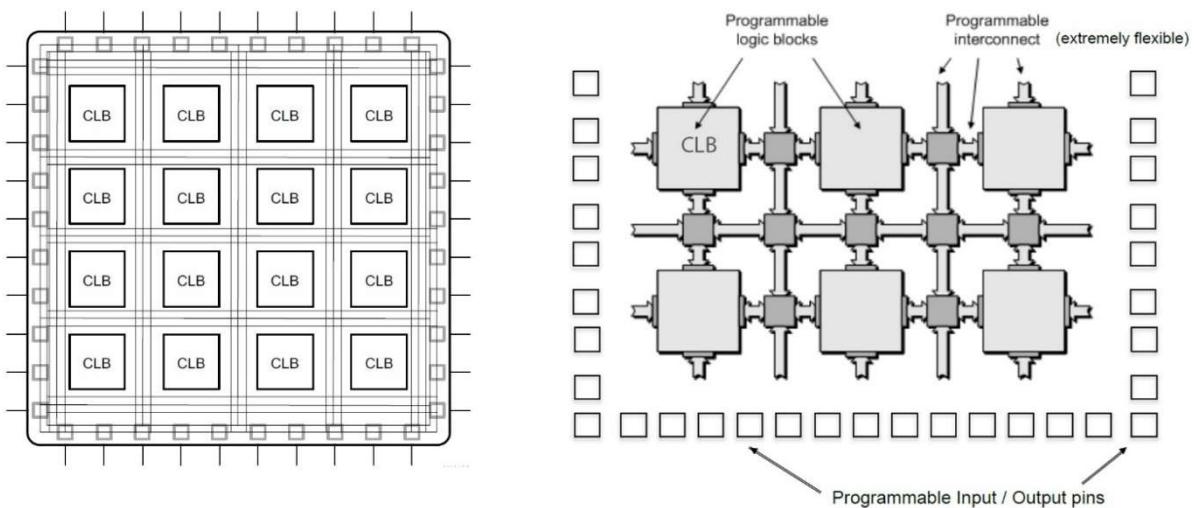


Figure 26 - Basic FPGA Architecture

The basic structure of an FPGA is composed of the following elements:

- **Look-up table (LUT):** This element performs logic operations.
- **Flip-Flop (FF):** This register element stores the result of the LUT.
- **Multiplexer:** a device that can receive multiple input signals and synthesize a single output signal.
- **Input/Output (I/O) pads:** These physically available ports get data in and out of the FPGA.

The combination of these elements results in the basic FPGA architecture shown in figure below.

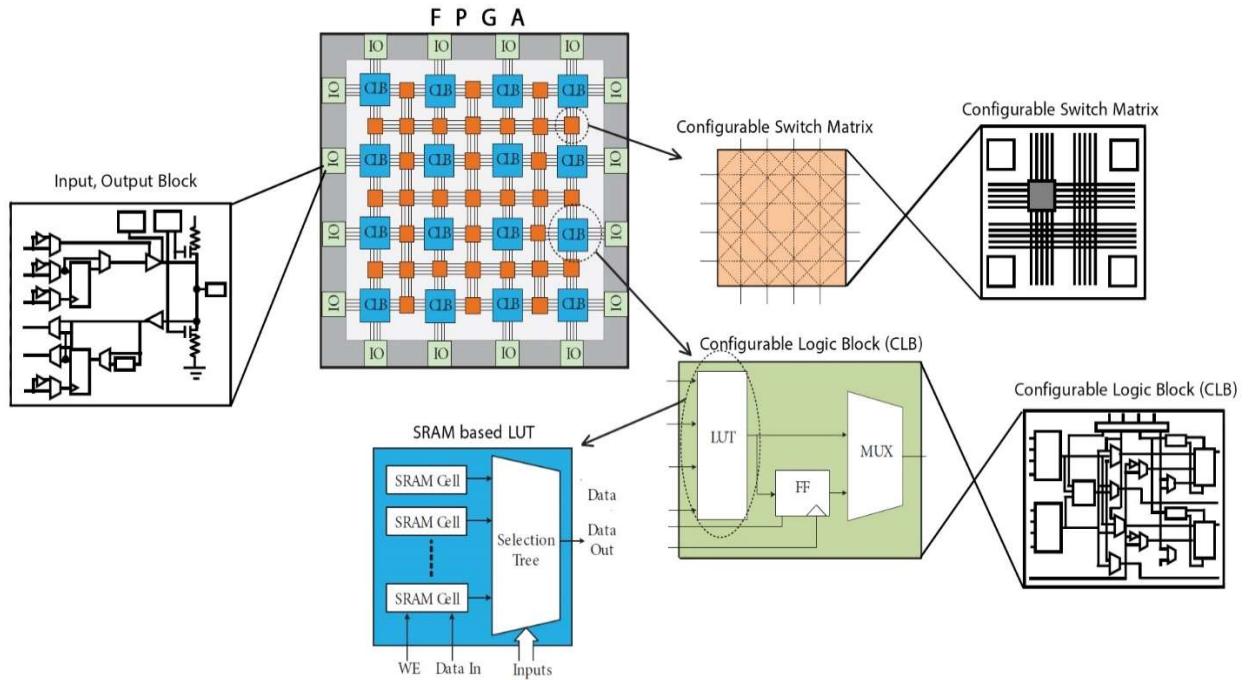


Figure 27 - FPGA Architecture

Internal Architecture of LUT

A combinational logic is made up of basic logic gate connected through wire. In FPGA these gates are simulated using look-up table (LUT). A n -input LUT requires 2^n bits of SRAM to store lookup table and $2^n - 1$ multiplexer to read individual bit. In general, 4-input look-up table is used but nowadays 6 input look-up table is used.

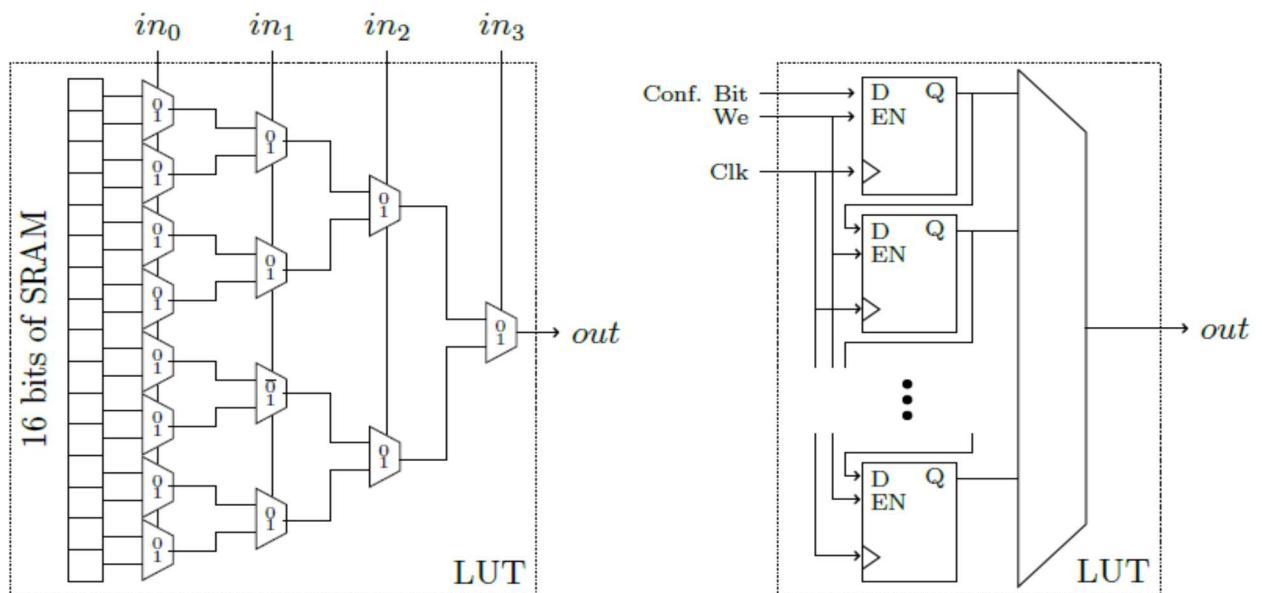


Figure 28 - Four input LUT - (left) circuitry for read - (right) circuitry for write.

The value in_0 to in_3 are used to determine which SRAM bit is given at output. Boolean function is stored in SRAM cell. SRAM cell is simply a shift register with one-bit width and 2^n bit depth. The Bit is shifted bit-by-bit into LUT when FPGA is programmed. LUT can also be used as memory element on

FPGA. When FPGA is programmed LUT is used as distributed RAM. The multiple LUTs are combined to make wider or deeper memories.

Routing Architecture - switch matrix

Interconnect is used for communicate between different logic island (LIs) these interconnects are configurable. It consists of horizontal and vertical channels (bundle of wire). At interconnect of routing channel there is a programmable links which determine how wire is connected, how input output is routed in particular logic island. All wires are connected to additional three wires at interconnection point but which connection is active is determine by programmable switch.

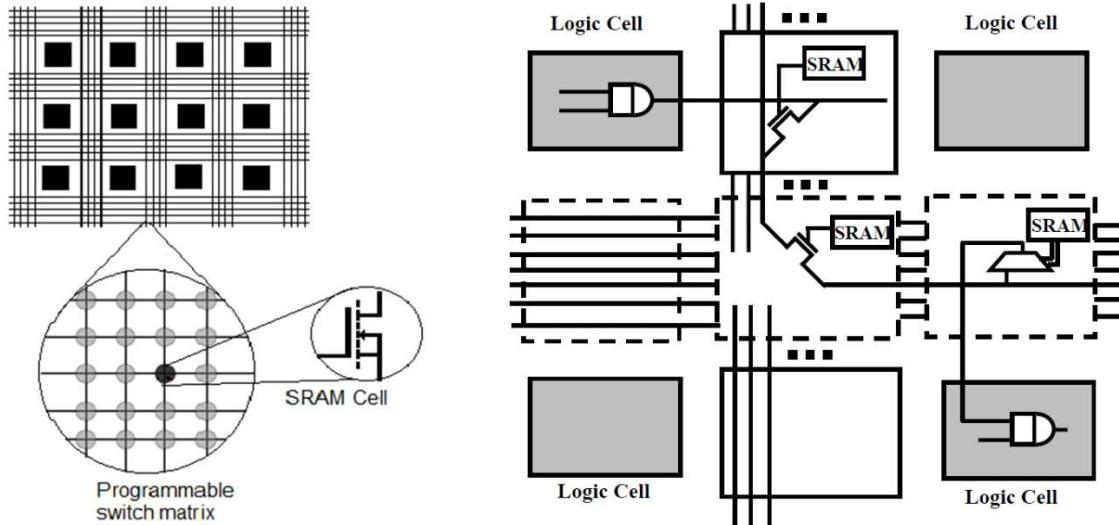


Figure 29 – (left) Routing architecture with switch matrix - (right) SRAM-controlled Programmable Switches

figure above, shows two applications of SRAM cells: for controlling the gate nodes of pass-transistor switches and to control the select lines of multiplexers that drive logic block inputs. The figure(right) gives an example of the connection of one logic block (represented by the AND-gate in the upper left corner) to another through two pass-transistor switches, and then a multiplexer, all controlled by SRAM cells.

Additional Elements In Contemporary FPGA Architectures

Contemporary FPGA architectures incorporate the basic elements along with additional computational and data storage blocks that increase the computational density and efficiency of the device. These additional elements are:

- Embedded memories for distributed data storage
- Phase-locked loops (PLLs) for driving the FPGA fabric at different clock rates
- High-speed serial transceivers
- Off-chip memory controllers
- Multiply-accumulate blocks (DSP Blocks¹²)

¹² The most complex computational block available in a common FPGA is the DSP block which is an arithmetic logic unit (ALU).

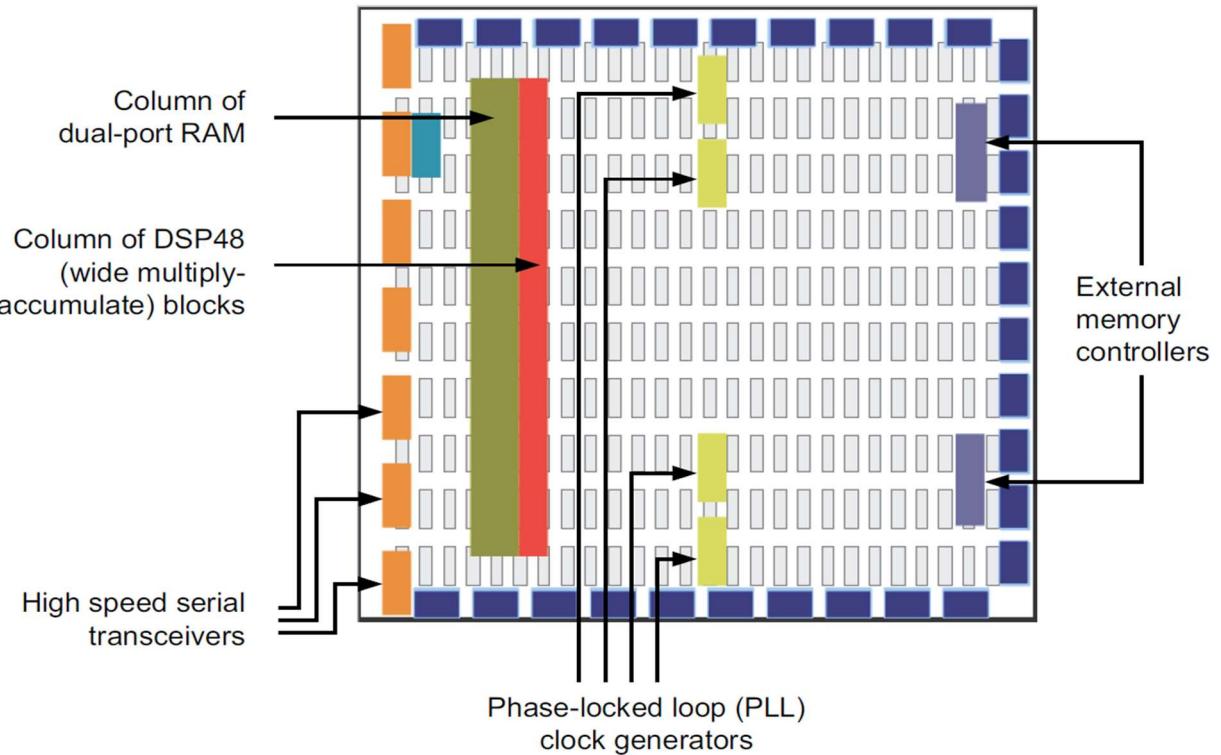


Figure 30 - Contemporary FPGA Architecture

The combination of these elements provides the FPGA with the flexibility to implement any software algorithm running on a processor.

4.3-FPGA Programming

The most important steps and tools of the design flow to produce an FPGA-circuit are depicted in figure below.

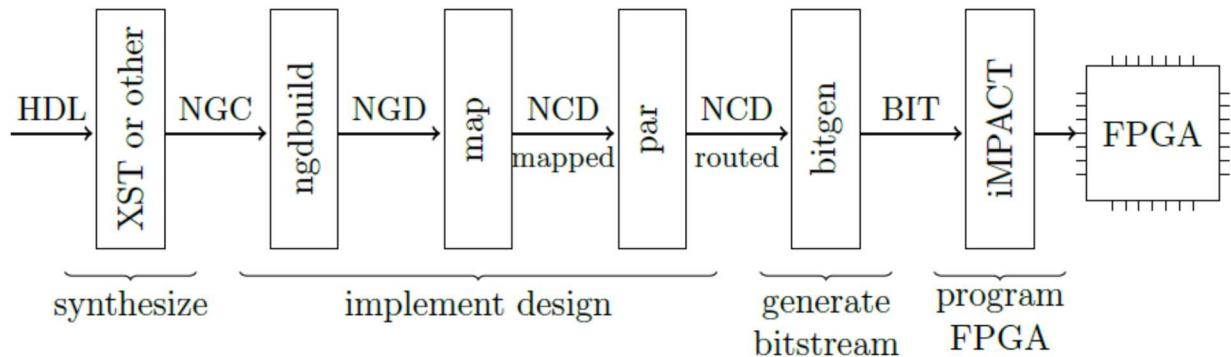


Figure 31 - FPGA design flow.

HDL & Synthesize

Programming on FPGA is same as describe the structure and behavior of electronic circuits and connecting wires in circuit, this is done by using a hardware description language (HDL) such as **VHDL** or **Verilog**. Synthesizer converts HDL into gate level netlist. A gate level netlist is basically your fitted design, before its converter to a programming file. It contains all of the logic and delays of the final system. It allows you to use your testbench from the simulation testing to test the final design. So, synthesis process generates netlist for each design element. The resulting netlists is saved to an **NGC (Native Generic Circuit)** file¹³.

Translate process

Translate process combines all the input netlists and constraints to a logic design file. This information is saved as a **NGD (Native Generic Database)** file. This can be done using NGD Build program. Here, defining manufacturing constraints is assigning the ports in the design to the physical elements (e.g., I/O pins, clock, switches, buttons etc.) Of the targeted device and specifying time requirements of the design. This information is stored in a file named **UCF (User Constraints File)**. Generally, NGC allow behavioral simulation and NGD allow timing simulation.

Map process

MAP process divides the whole circuit with logical elements into sub blocks such that they can be fit into the FPGA logic blocks. That means map process fits the logic defined by the NGD file into the targeted FPGA elements (configurable Logic Blocks (CLB), Input Output Blocks (IOB)) and generates an **NCD (Native Circuit Description)** file which physically represents the design mapped to the components of FPGA.

Place and Route

PAR program is used for this process. The place and route process places the sub blocks from the map process into logic blocks according to the constraints and connects the logic blocks. A tradeoff between all the constraints is taken account by the place and route process. The PAR tool takes the

¹³ For Xilinx Synthesis Technology (XST)

mapped NCD file as input and produces a completely routed NCD file as output. Output NCD file consists of the routing information.

Device Programming - Bitstream Generation

Now the design must be loaded on the FPGA. But the design must be converted to a format so that the FPGA can accept it. BITGEN program deals with the conversion. It encodes design in binary known as **Bitstream (a .BIT file)**. The routed NCD file is given to the BITGEN program to generate a bitstream which can be used to configure the target FPGA device. Inside FPGA a finite state machine control by Bitstream which extract configuration data from Bitstream. This can be done using a cable.

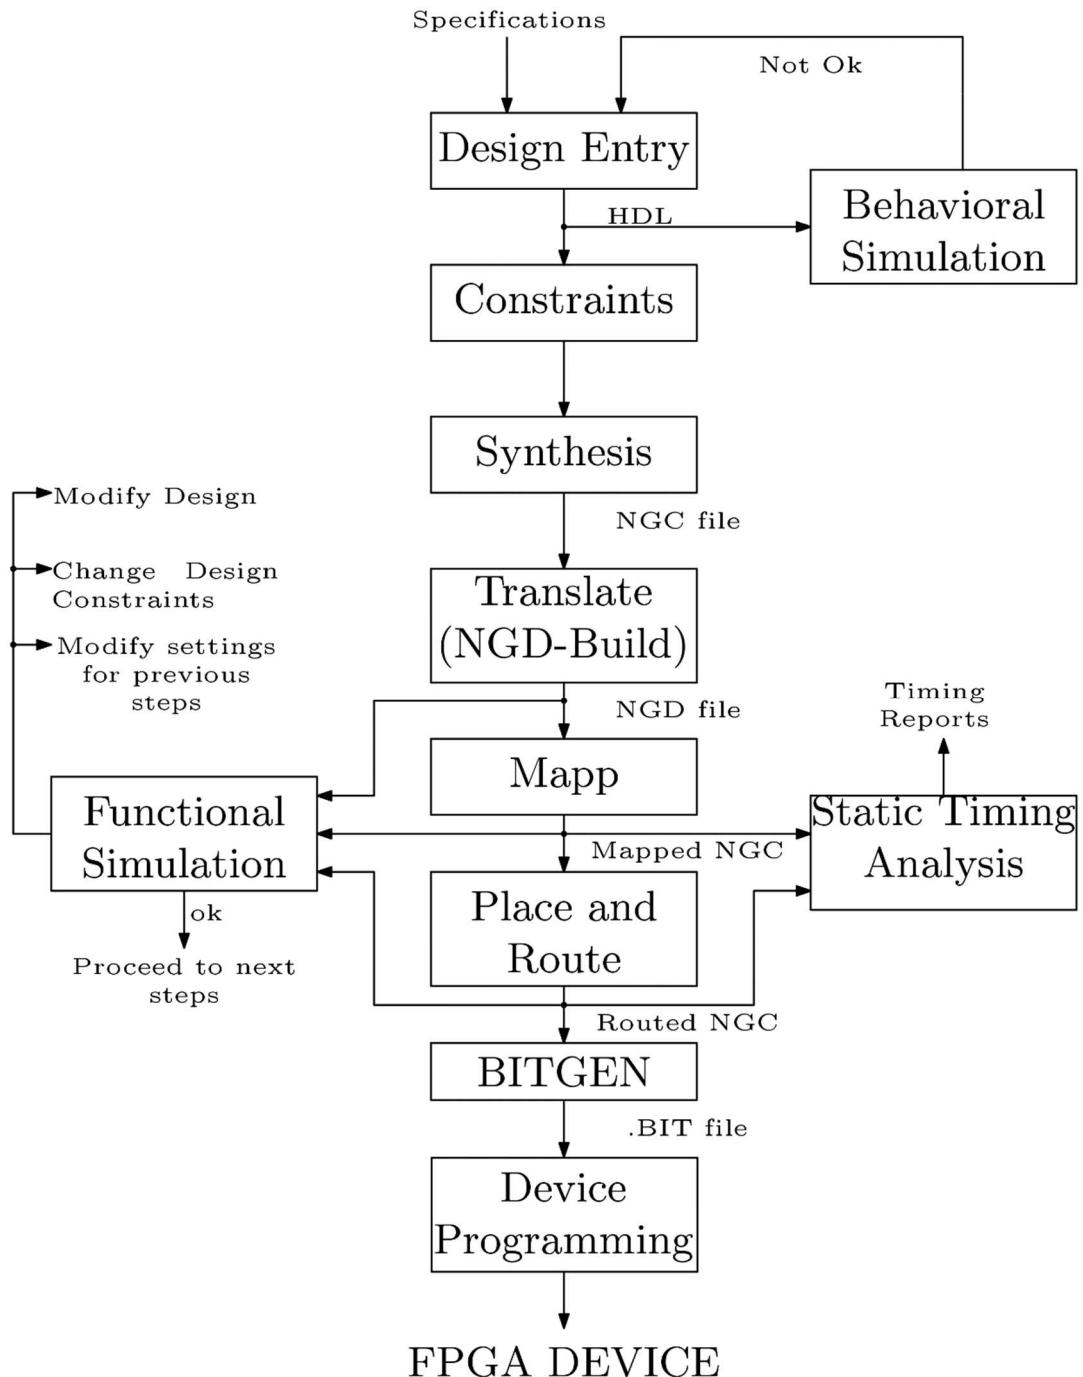


Figure 32 - FPGA implementation - Step by Step

4.4-FPGAs Vs ASICs

ASIC microchip is designed and manufactured for one specific application and does not allow you to reprogram or modify it after it is produced. This means ASICs are not intended for general use. For chose technology to design or develop a chip, there are several factors to consider:

- Design Flow
- Flexibility
- Performance and Efficiency
- Cost and Time to Market
- Power Consumption

The main properties and applications of ASICs and FPGAs are listed below:

FPGAs

- Off-the-shelf, that is available immediately and does not need to be specially made to suit a particular purpose.
- Reconfigurability, may be reprogrammed in the field (gateware upgrade) – (New features - Bug fixes)
- Rapid development cycle (minutes / hours) - Short time to market
- Low development cost

ASICs

- Higher performance - Low power
- Analog designs possible
- Long development cycle (weeks /months)
- Design cannot be changed once it is produced
- Better radiation hardness
- Extremely high development cost (ASICs are produced at a semiconductor fabrication facility ("fab") according to your design)
- Lower cost per device compared to FPGA, when large quantities are needed

	ASIC	FPGA
Design Flow	Complex	Simplistic
Flexibility		✓
Performance	✓	
Development Cost (NRE)		✓
Production Unit Cost	✓	
Power Consumption	✓	

Figure 33 - FPGAs Vs ASICs

4.5-FPGA As Bitcoin Miner

5-CONCLUSION

Bitcoin is a new Internet currency that anyone can get started mining. There are a number of reasons you might mine: for profit, to help secure the network, to help found a new Internet currency, or just to gain technical experience.

Professional mining

Today mining has mostly moved away from individuals and toward professional mining centers. Exact details about how these centers operate are not very well known because companies want to protect their setups to maintain a competitive advantage. Presumably, these operations maintain profitability by buying slightly newer and more efficient ASICs than are available for general sale at a bulk discount. Figure below shows a picture of a professional mining center.



Figure 34 - BitFury mining center, a professional mining center in the republic of Georgia

Mining hardware evolution: Similarities to gold mining

Currently, ASIC mining is the only realistic means to be profitable in Bitcoin and it's not very friendly to small miners. Whereas with Bitcoin mining we've seen an evolution from CPUs to GPUs to FPGAs, to now ASICs, gold mining saw an evolution from individuals with gold pans to small groups of people with sluice boxes, to placer mining¹⁴ to modern gold mining which often utilizes gigantic open pit mines to extract tons of raw material from the earth (See Figure below).

¹⁴ consisting of large mining groups blasting away hillsides with water

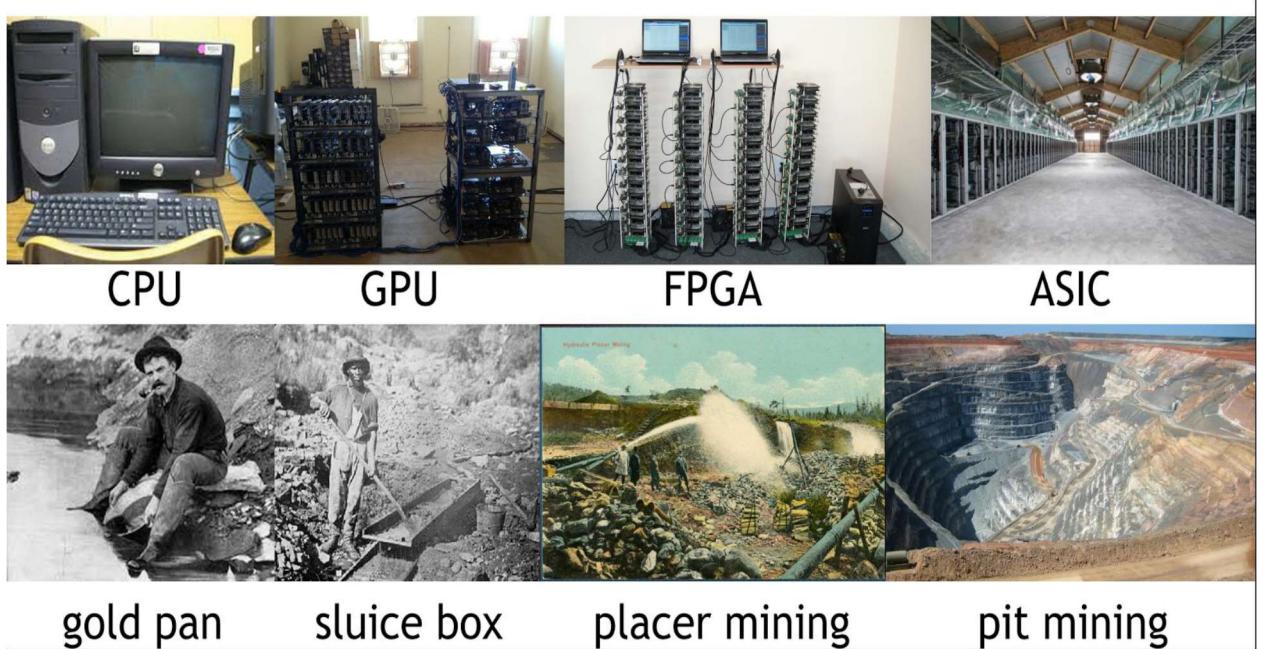


Figure 35 - Evolution of mining. We can see a clear parallel between the evolution of Bitcoin mining and the evolution of gold mining. Both were initially friendly to individuals and over time became massive operations controlled by large companies.

Both with Bitcoin and with gold, the friendliness and accessibility to individuals has gone down over time and large companies have eventually consolidated most of the operations (and profits).